

11Oct2022

Day06

## Minikube installation

### Kubernetes Workloads I

#### Install Minikube all-in-one

<https://minikube.sigs.k8s.io/docs/start/>

#### Minikube Installation Prerequisites

2 vCPUs

2GB memory

30GB HDD

Internet connection

#### Step 1. create a VM

download ubuntu iso image

<https://releases.ubuntu.com/18.04/ubuntu-18.04.6-live-server-amd64.iso>

<https://releases.ubuntu.com/focal/ubuntu-20.04.5-live-server-amd64.iso>

create template and install ubuntu linux through iso image

#### config ubuntu VM

login with **devops** user

u: devops

p: ubuntu

\$ sudo -i

password: **ubuntu**

# passwd root

set pass: **ubuntu**

# hostnamectl set-hostname **minikube.example.com**

# ifconfig

# route -n

# cat /etc/resolv.conf

# vim /etc/netplan/00-installer-config.yaml

network:

version: 2

**renderer: networkd**

ethernets:

enp0s3:

dhcp4: **false**

**addresses: [192.168.29.34/24]**

**gateway4: 192.168.29.1**

**nameservers:**

**addresses: [8.8.8.8,8.8.4.4,192.168.29.1]**

:wq!

# netplan apply

# vim /etc/ssh/sshd\_config

33 PermitRootLogin yes

:wq!

# systemctl restart sshd.service

# poweroff

**NOTE:** take snapshot

#### Step 2. Install Kubectl

Kubectl is the command line tool for Kubernetes to interact with the cluster from the command line.

# apt-get update && sudo apt-get install -y apt-transport-https gnupg2

# curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -

# echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee -a /etc/apt/sources.list.d/kubernetes.list

# apt-get update

# apt-get install -y kubectl

# kubectl version

#### Step 3. Install additional dependencies

# apt-get install conntrack

#### Step 4. Install minikube

# apt-get install -y docker.io

# curl -Lo minikube <https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64>

# chmod +x minikube

# mv minikube /usr/local/bin

# set NO\_PROXY=localhost,127.0.0.1,192.168.29.0/24

## Docker, Kubernetes, VMware Tanzu and RedHat OCP

### Step 5. Install cri-dockerd

5-1. wget the file

```
# wget https://github.com/Mirantis/cri-dockerd/releases/download/v0.2.0/cri-dockerd-v0.2.0-linux-amd64.tar.gz
```

5-2. unzip the package. On Linux you can use

```
# tar xvf cri-dockerd-v0.2.0-linux-amd64.tar.gz
```

5-3. move the cri-dockerd binary to your usr/local/bin directory

```
# mv ./cri-dockerd /usr/local/bin/
```

5-4. check if it is successfully installed

```
# cri-dockerd --help
```

5-5. start the service on Linux

```
# wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-docker.service
```

```
# wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-docker.socket
```

```
# mv cri-docker.socket cri-docker.service /etc/systemd/system/
```

```
# sed -i -e 's,/usr/bin/cri-dockerd,/usr/local/bin/cri-dockerd,' /etc/systemd/system/cri-docker.service
```

```
# systemctl daemon-reload
```

```
# systemctl enable cri-docker.service
```

```
# systemctl enable --now cri-docker.socket
```

```
# systemctl status cri-docker.socket
```

### Step 6. Install CRI-O Container Runtime on Ubuntu

```
# apt update && sudo apt upgrade
```

```
# OS=xUbuntu_18.04
```

```
# CRI_O_VERSION=1.23
```

```
# echo "deb https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/$OS/ /" | sudo tee /etc/apt/sources.list.d/devel:kubic:libcontainers:stable.list
```

```
# echo "deb http://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable:/cri-o:/$CRI_O_VERSION/$OS/ /" | sudo tee /etc/apt/sources.list.d/devel:kubic:libcontainers:stable:cri-o:$CRI_O_VERSION.list
```

```
# curl -L https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable:cri-o:$CRI_O_VERSION/$OS/Release.key | sudo apt-key add -
```

```
# curl -L https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/$OS/Release.key | sudo apt-key add -
```

```
# apt update
```

```
# apt install cri-o cri-o-runc
```

```
# systemctl enable cri-o.service
```

```
# systemctl start cri-o.service
```

```
# apt install cri-tools
```

### Step 7. Start the minikube

```
# minikube start --vm-driver=none --docker-env NO_PROXY=$NO_PROXY
```

#### after install success

```
# sudo mv /root/.kube /root/.minikube $HOME
```

```
# sudo chown -R $USER $HOME/.kube $HOME/.minikube
```

```
# minikube status
```

```
minikube
```

```
type: Control Plane
```

```
host: Running
```

```
kubelet: Running
```

```
apiserver: Running
```

```
kubeconfig: Configured
```

```
# minikube completion bash >> ~/.bashrc
```

```
# kubectl completion bash >> ~/.bashrc
```

```
# source ~/.bashrc
```

```
# echo "autocmd FileType yaml setlocal ai ts=2 sw=2 et cursorcolumn" >> ~/.vimrc
```

```
# kubectl get all -A
```

#### References

<https://www.mirantis.com/blog/how-to-install-cri-dockerd-and-migrate-nodes-from-dockershim/>

<https://developer.ibm.com/tutorials/set-up-minikube-on-ubuntu-server-within-minutes/>

<https://computingforgeeks.com/install-cri-o-container-runtime-on-ubuntu-linux/>

## Workloads

<https://kubernetes.io/docs/concepts/workloads/>

A workload is an application running on Kubernetes.

1-Pod

## Pods

<https://kubernetes.io/docs/concepts/workloads/pods/>

Pods are the smallest deployable units of computing that you can create and manage in Kubernetes.

### Create

#### cli

```
# kubectl run <pod-name> --image <image-name>
```

```
# kubectl run pod1 --image nginx --dry-run=client
```

```
# kubectl run pod1 --image nginx
```

#### file

```
# kubectl run pod1 --image nginx -o yaml --dry-run=client >/tmp/pod1.yaml
```

```
# kubectl run pod1 --image nginx -o json --dry-run=client >/tmp/pod1.json
```

```
# cat /tmp/pod1.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  labels:
```

```
    run: pod1
```

```
  name: pod1
```

```
spec:
```

```
  containers:
```

```
    - image: nginx
```

```
      name: nginxcnt
```

```
# kubectl create -f /tmp/pod1.yaml --dry-run=client
```

```
# kubectl create -f /tmp/pod1.yaml
```

### operate

```
# kubectl get pod
```

```
# kubectl describe pod pod1
```

```
# kubectl edit pod pod1
```

```
# kubectl delete pod pod1 --force --grace-period=0
```

### edit pod parameters

#### upgrade/downgrade the image

##### online

```
# kubectl get pod pod1 -o yaml | grep -i "image"
```

```
  - image: nginx
```

```
# kubectl edit pod pod1
```

```
19 spec:
```

```
20   containers:
```

```
21     - image: nginx:1.18
```

```
:wq!
```

```
# kubectl get pod pod1 -o yaml | grep -i "image"
```

```
  - image: nginx:1.18
```

##### offline

```
# kubectl get pod pod1 -o yaml >/tmp/pod1running.yaml
```

```
# vim /tmp/pod1running.yaml
```

```
# kubectl replace -f /tmp/pod1running.yaml
```

```
# kubectl get pod pod1 -o yaml | grep -i "image"
```

```
  - image: nginx:1.19
```

create and replace      apply and apply

### Backup

```
kubectl get pod pod1 -o yaml >/tmp/pod1back.yaml
```

### Delete

```
# kubectl delete pod pod1 --force --grace-period=0
```

### Restore

```
# kubectl apply -f /tmp/pod1back.yaml
```

### Log

```
# kubectl logs pods/pod1
```

### get inside pod

```
# kubectl exec -it pod/pod1 -- /bin/bash
```

```
root@pod1:/# exit
```

## Docker, Kubernetes, VMware Tanzu and RedHat OCP

exit

or

# kubectl exec -it pod/pod1 -- ls /

### run workloads in different namespace

#### 1-change namespace and create resource

# kubectl config get-contexts

# kubectl config set-context --namespace test1 --current

# kubectl config get-contexts

\* kubernetes-admin@kubernetes kubernetes kubernetes-admin test1

# kubectl run pod1 --image redis

# kubectl get pods

pod1 1/1 Running 0 6s

#### 2-create resource without namespace

# kubectl config get-contexts

\* kubernetes-admin@kubernetes kubernetes kubernetes-admin test1

# kubectl run redis --image redis --namespace <namespace>

# kubectl run redis --image redis --namespace default

# kubectl get pods --namespace default

or

# kubectl run pod3 --image alpine -o yaml --dry-run=client >alpine.yaml

# vim alpine.yaml

cat alpine.yaml

apiVersion: v1

kind: Pod

metadata:

labels:

run: pod3

name: pod3

namespace: test1

spec:

containers:

- image: alpine

name: alpinecnt

:wq!

# kubectl apply -f alpine.yaml

# kubectl get pods --namespace test1

#### create Pod with multiple containers

# kubectl run pod1 --image nginx -o yaml --dry-run=client >pod1.yaml

# vim pod1.yaml

apiVersion: v1

kind: Pod

metadata:

labels:

run: pod1

name: pod1

spec:

containers:

- image: nginx

name: nginxcnt

- image: redis

name: rediscnt

:wq!

# kubectl create -f pod1.yaml --dry-run=client

# kubectl create -f pod1.yaml

# kubectl get pods -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod1	2/2	Running	0	106s	192.168.1.3	node01

#### get inside Redis Container

# kubectl exec -it pod/pod1 -c redis -- /bin/bash

root@pod1:/data#

exit

-c container

#### get nginx container Log

# kubectl logs pod/pod1 -c nginxcnt

# kubectl logs pod/pod1 -c redis

#### send/receive file

kubectl cp /tmp/f1 pod1:/tmp -c nginxcnt

kubectl exec -it pod/pod1 -c nginxcnt -- ls /tmp

