

## Workloads

<https://kubernetes.io/docs/concepts/workloads/>

A workload is an application running on Kubernetes.

1-Pod

2-ReplicaSet

3-Deployment

5-DaemonSet

6-Init Containers

7-Static Pod

## DaemonSet

<https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

A DaemonSet ensures that all Nodes run a copy of a Pod.

As nodes are added to the cluster, Pods are added to them.

As nodes are removed from the cluster, those Pods are garbage collected. Deleting a DaemonSet will clean up the Pods it created.

ex

```
# kubectl get nodes
```

```
# kubectl api-resources | grep -i "daemon"
```

```
daemonsets      ds      apps/v1      true      DaemonSet
```

```
# vim rsyslogds.yaml
```

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  name: rsyslogds
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      name: rsyslog
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        name: rsyslog
```

```
    spec:
```

```
      containers:
```

```
        - name: rsyslogcnt
```

```
          image: rsyslog/syslog_appliance_alpine      ->selected image from specific registry
```

```
:wq!
```

```
# kubectl apply -f rsyslogds.yaml --dry-run=client
```

```
# kubectl apply -f rsyslogds.yaml
```

```
# watch kubectl get pods -o wide
```

```
# kubectl get ds
```

```
# kubectl describe ds rsyslogds
```

## Init Containers

<https://kubernetes.io/docs/concepts/workloads/pods/init-containers/>

-initContainers: specialized containers that run before app containers in a Pod.

-initContainers can contain utilities or setup scripts not present in an app image.

-initContainers, which are run before the app containers are started.

-each initContainer must complete successfully before the next one will start.

ex

```
# vim initCnt.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx-pod
```

```
spec:
```

```
  initContainers:
```

```
    - name: busyboxcnt
```

```
      image: busybox
```

```
  containers:
```

```
    - name: nginxcnt
```

```
      image: nginx
```

```
:wq!
```

```
# kubectl apply -f initCnt.yaml
```

```
# watch kubectl get pods
```

**static Pods**

<https://kubernetes.io/docs/tasks/configure-pod-container/static-pod/>

-Static Pods are managed directly by the kubelet daemon on a specific node, without the API server observing them.

-Unlike Pods that are managed by the control plane; instead, the kubelet watches each static Pod.

-default location for static pod is **/etc/kubernetes/manifests/**

**How change StaticPod default location**

in MasterNode/ControlPlane edit:

```
# vim /var/lib/kubelet/config.yaml
```

```
41 staticPodPath: /etc/kubernetes/manifests
```

->this path should be change

```
:wq!
```

**Now**, StaticPods will store in new location

**ex**

**create StaticPod on Node3**

```
root@master1:~# hostname
```

```
master1.example.com
```

```
root@master1:~# kubectl get node -o wide
```

```
root@master1:~# ssh root@node3.example.com
```

```
root@node3:~# hostname
```

```
node3.example.com
```

```
root@node3:~# echo "autocmd FileType yaml setlocal ai ts=2 sw=2 et cursorcolumn" > ~/.vimrc
```

```
root@node3:~# cd /etc/kubernetes/manifests/
```

```
root@node3:/etc/kubernetes/manifests# pwd
```

```
/etc/kubernetes/manifests
```

```
root@node3:/etc/kubernetes/manifests# vim redis.yaml
```

```
kind: Pod
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: redispod
```

```
spec:
```

```
  containers:
```

```
    - name: rediscont
```

```
      image: redis
```

```
:wq!
```

```
# ls
```

```
redis.yaml
```

```
root@node3:~# exit
```

```
root@master1:~# kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
redispod-node3.example.com	1/1	Running	0	87s	10.85.0.5	node3.example.com

**How assign Pod to specific Node**

1-nodeSelector field

2-nodeName field

3-Affinity and anti-affinity

**1-nodeSelector field**

```
# kubectl get nodes node2.example.com --show-labels
```

```
# kubectl label nodes node2.example.com lbl=node
```

```
# kubectl run redispod --image redis -o yaml --dry-run=client > redispod.yaml
```

```
# vim redispod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: redispod
```

```
spec:
```

```
  containers:
```

```
    - image: redis
```

```
      name: rediscont
```

```
nodeSelector:
```

```
  lbl: node
```

```
:wq!
```

```
# kubectl create -f redispod.yaml
```

```
# kubectl get pods -o wide
```

redispod	1/1	Running	0	79s	10.85.0.5	node2.example.com
----------	-----	---------	---	-----	-----------	-------------------

**2-nodeName field**

```
# kubectl run nginx-pod --image nginx -o yaml --dry-run=client > nginxpod.yaml
```

```
# vim nginxpod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginxpod
```

```
spec:
```

```
  containers:
```

```
    - image: nginx
```

```
      name: nginxcnt
```

```
nodeName: node4.example.com
```

```
:wq!
```

```
# kubectl apply -f nginxpod.yaml --dry-run=client
```

```
# kubectl apply -f nginxpod.yaml
```

```
# kubectl get pods -o wide
```

nginxpod	1/1	Running	0	51s	10.85.0.3	node4.example.com
----------	-----	---------	---	-----	-----------	-------------------

## Attach specific namespace to specific Node/Nodes

**step1-** enable **PodNodeSelector** on kube-api component

**step2-** create namespace

**step3-** label Node and Namespace

### step1- Create Namespace

```
# kubectl create namespace test1 -o yaml --dry-run=client >test1.yaml
# ls
test1.yaml
# cat test1.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: test1
:wq!
# kubectl apply -f test1.yaml
# kubectl get ns
```

**step2-** enable **PodNodeSelector** on kube-api component

```
# vim /etc/kubernetes/manifests/kube-apiserver.yaml
20  --enable-admission-plugins=PodNodeSelector
:wq!
# watch kubectl get nodes
```

### step3- label Node and Namespace

**add**

```
# kubectl label nodes node5.example.com ns=test1
# kubectl get nodes node5.example.com --show-labels
```

**remove**

```
# kubectl label nodes node5.example.com ns-
# kubectl get nodes node5.example.com --show-labels
```

**update label on Namespace**

```
# vim test1.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: test1
  annotations:
    scheduler.alpha.kubernetes.io/node-selector: ns=test1
:wq!
# kubectl apply -f test1.yaml
verify
# kubectl run nginx --image nginx --namespace test1
# watch kubectl get pods --namespace test1 -o wide
nginx 1/1 Running 0 82s 10.85.0.3 node5.example.com
```