

15Oct2022

Day09

Kubernetes Workloads V

Kubernetes Resource Manager

Workloads

<https://kubernetes.io/docs/concepts/workloads/>

A workload is an application running on Kubernetes.

1-Pod

2-ReplicaSet

3-Deployment

5-DaemonSet

6-Init Containers

7-Static Pod

8-Job

9-CronJob

Jobs

<https://kubernetes.io/docs/concepts/workloads/controllers/job/>

-A Job creates one or more Pods and will continue to retry execution of the Pods until a specified number of them successfully terminate. As pods successfully complete, the Job tracks the successful completions. When a specified number of successful completions is reached, the task is complete.

-Deleting a Job will clean up the Pods it created.

ex

```
# kubectl api-resources | grep -i "job"
```

```
jobs                batch/v1            true             Job
```

```
# vim alpinejob.yaml
```

```
kind: Job
```

```
apiVersion: batch/v1
```

```
metadata:
```

```
  name: alpinejob
```

```
spec:
```

```
  template:
```

```
    spec:
```

```
      containers:
```

```
        - name: alpinecnt
```

```
          image: alpine
```

```
          command: ["/bin/sh", "-c"]
```

```
          args: ["echo 'Hellooooo!';sleep 15"]
```

```
          restartPolicy: Never
```

```
:wq!
```

```
# kubectl apply -f alpinejob.yaml
```

```
# watch kubectl get job
```

```
# kubectl get pods -o wide
```

```
# kubectl logs alpinejob-sbjt
```

Parallel execution for Jobs

1-Controlling parallelism

2-Completion mode

1-Controlling parallelism

-all pods in job will start at the same time with enable option '**parallelism**'

-The requested parallelism can be set to any non-negative value.

ex

```
# vim alpinejob.yaml
```

```
kind: Job
```

```
apiVersion: batch/v1
```

```
metadata:
```

```
  name: alpinejob
```

```
spec:
```

```
  parallelism: 4
```

```
  template:
```

```
    spec:
```

```
      containers:
```

```
        - name: alpinecnt
```

```
          image: alpine
```

```
          command: ["/bin/sh", "-c"]
```

```
          args: ["echo 'Hellooooo!';sleep 15"]
```

```
          restartPolicy: Never
```

```
:wq!
```

```
# kubectl apply -f alpinejob.yaml
```

```
# watch kubectl get pods
```

```
# watch kubectl get jobs
```


Resource Management for Pods and Containers

<https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/>

-When you specify a Pod, you can optionally specify how much of each resource a container need.

-The most common resources to specify are:

- 1-CPU
- 2-memory (RAM)

Resource types

CPU and memory are each a resource type

A resource type has a base unit.

CPU represents compute processing and is specified in units of Kubernetes CPUs.

CPU measurement is in millicores. with this measure format each **1CPU is splinted into 1000units(millicores)**

Memory is specified in units of **bytes**.

Memory can express as plain integer in B, K, M, G, T, P

Monitoring resources through Metrix-Server

<https://kubernetes.io/docs/tasks/debug/debug-cluster/resource-metrics-pipeline/#metrics-server>

<https://github.com/kubernetes-sigs/metrics-server>

Metrics Server is a scalable, efficient source of container resource metrics for Kubernetes built-in autoscaling pipelines.

```
# kubectl top nodes
```

```
error: Metrics API not available
```

```
# kubectl top pod
```

```
error: Metrics API not available
```

Installation

<https://github.com/kubernetes-sigs/metrics-server>

```
# kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

```
# kubectl get deployments.apps -A
```

```
# kubectl get deployments.apps --namespace kube-system
```

```
# kubectl top nodes
```

```
Error from server (ServiceUnavailable): the server is currently unable to handle the request (get nodes.metrics.k8s.io)
```

```
# kubectl get deployments.apps --namespace kube-system
```

```
# kubectl get pods --namespace kube-system
```

```
# kubectl logs pods/metrics-server-678f4bf65b-jnm2f --namespace kube-system
```

to disable tls certification request,

```
# kubectl edit deployments.apps --namespace kube-system metrics-server
```

```
42 - --kubelet-preferred-address-types=InternalIP ->modify to 'InternalIP' remove others
```

```
43 - --kubelet-insecure-tls ->append this line
```

```
:wq!
```

```
# kubectl top nodes
```

```
Error from server (ServiceUnavailable): the server is currently unable to handle the request (get nodes.metrics.k8s.io)
```

```
# kubectl rollout restart deployment --namespace kube-system metrics-server
```

```
# kubectl top nodes
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
master1.example.com	152m	7%	1684Mi	43%
node1.example.com	60m	6%	543Mi	28%

Requests and Limits

Requests ->minimum requirement to start

when you specify the resource request for containers in a Pod, the kube-scheduler uses this information to decide which node to place the Pod on.

Limits ->maximum access to resources

when you specify a resource limit for a container, the kubelet enforces those limits so that the running container is not allowed to use more of that resource than the limit you set.

ex

```
# kubectl run nginx --image nginx -o yaml --dry-run=client >nginxpod.yaml
```

```
# vim nginxpod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

```
spec:
```

```
  containers:
```

```
  - image: nginx
```

```
    name: nginx
```

```
    resources:
```

```
      requests: ->to start, minimum
```

```
        cpu: "4m"
```

```
        memory: "2Mi"
```

```
      limits: ->to make limit, maximum
```

```
        cpu: "8m"
```

```
        memory: "4Mi"
```

```
:wq!
```

```
# kubectl top pods
```

NAME	CPU(cores)	MEMORY(bytes)
redisj-27764091-tfq6k	2m	2Mi