

17Oct2022

Day10

K8s MultiNode Cluster Installation**Prepare Hosts to install K8s MultiNode Cluster****1MasterNode**

- A compatible Linux host
- 4GB or more of RAM per machine.
- 2CPUs or more.
- Full network connectivity between all machines in the cluster (public or private network is fine).
- Unique hostname, MAC address, and product_uuid for every node.

2WorkerNode

- A compatible Linux host
- 2GB or more of RAM per machine.
- 1CPUs or more.
- Full network connectivity between all machines in the cluster (public or private network is fine).
- Unique hostname, MAC address, and product_uuid for every node.

after finish Installation power off host and make clone.

master1.example.com

worker1.example.com

worker2.example.com

MasterNode

```
# hostnamectl set-hostname k8smaster1.example.com
```

```
# ip a s
```

```
# nmcli connection show
```

```
NAME UUID TYPE DEVICE
ens192 5bb5771d-b4df-4997-8d59-0e5c41fdb5b1 ethernet ens192
```

```
# nmcli connection delete "ens192"
```

```
# nmcli connection reload
```

```
# nmcli connection add con-name "ens192" type ethernet autoconnect yes ifname ens192 ipv4.address "192.168.0.191/24" ipv4.dns "192.168.0.1" ipv4.gateway "192.168.0.1" +ipv4.dns "8.8.8.8" +ipv4.dns "8.8.4.4" ipv4.method manual
```

```
# nmcli connection reload
```

```
# vi /etc/hosts
```

```
192.168.0.191 k8smaster1.example.com k8smaster1
```

```
192.168.0.192 k8sworker1.example.com k8sworker1
```

```
192.168.0.193 k8sworker2.example.com k8sworker2
```

```
192.168.0.194 k8sworker3.example.com k8sworker3
```

```
:wq!
```

```
# lsblk
```

```
sr0 11:0 110.7G 0 rom /media/cdrom
```

```
# mkdir /media/cdrom
```

```
# echo "/dev/sr0 /media/cdrom iso9660 defaults 0 0" >>/etc/fstab
```

```
# mount -a
```

```
# vi /etc/yum.repos.d/redhat.repo
```

```
[App]
```

```
name=AppStream
```

```
baseurl=file:///media/cdrom/AppStream/
```

```
gpgcheck=0
```

```
enabled=1
```

```
[BaseOS]
```

```
name=BaseOS
```

```
baseurl=file:///media/cdrom/BaseOS/
```

```
gpgcheck=0
```

```
enabled=1
```

```
:wq!
```

```
# yum repolist
```

```
poweroff, take snapshot
```

WorkerNodes

```
hostname k8sworker1.example.com
```

```
k8sworker2.example.com
```

```
k8sworker3.example.com
```

Config Network

```
192.168.0.192/24
```

```
192.168.0.193/24
```

```
192.168.0.194/24
```

Config Local DNS

```
192.168.0.191 k8smaster1.example.com k8smaster1
```

```
192.168.0.192 k8sworker1.example.com k8sworker1
```

```
192.168.0.193 k8sworker2.example.com k8sworker2
```

```
192.168.0.194 k8sworker3.example.com k8sworker3
```

Config Local Repository

```
/etc/yum.repos.d/rhel.repo
```

```
poweroff, make snapshot
```

Install K8s MultiNode Cluster

<https://kubernetes.io/>

<https://kubernetes.io/docs/home/>

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

Master/Workers

Installing a container runtime

To run containers in Pods, Kubernetes uses a container runtime.

Container Runtimes

Option 1: From the official binaries

need to install a container runtime into each node in the cluster so that Pods can run there.

Kubernetes 1.25 requires that you use a runtime that conforms with the **Container Runtime Interface-CRI**

1- containerd

<https://github.com/containerd/containerd/blob/main/docs/getting-started.md>

<https://github.com/containerd/containerd/releases>

1-1. Installing containerd

```
# yum install wget -y
```

```
# wget https://github.com/containerd/containerd/releases/download/v1.6.8/containerd-1.6.8-linux-amd64.tar.gz
```

```
# ls containerd-1.6.8-linux-amd64.tar.gz
```

```
# tar Cxvf /usr/local containerd-1.6.8-linux-amd64.tar.gz 1-
```

2. systemd

```
# mkdir -p /usr/local/lib/systemd/system/
```

```
# vi /usr/local/lib/systemd/system/containerd.service
```

```
# Copyright The containerd Authors.
```

```
#
```

```
# Licensed under the Apache License, Version 2.0 (the "License");
```

```
# you may not use this file except in compliance with the License.
```

```
# You may obtain a copy of the License at
```

```
#
```

```
# http://www.apache.org/licenses/LICENSE-2.0
```

```
#
```

```
# Unless required by applicable law or agreed to in writing, software
```

```
# distributed under the License is distributed on an "AS IS" BASIS,
```

```
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
```

```
# See the License for the specific language governing permissions and
```

```
# limitations under the License.
```

```
[Unit]
```

```
Description=containerd container runtime
```

```
Documentation=https://containerd.io
```

```
After=network.target local-fs.target
```

```
[Service]
```

```
#uncomment to enable the experimental sb service (sandboxed) version of containerd/cni integration
```

```
#Environment="ENABLE_CRI_SANDBOXES=sandboxed"
```

```
ExecStartPre=-/sbin/modprobe overlay
```

```
ExecStart=/usr/local/bin/containerd
```

```
Type=notify
```

```
Delegate=yes
```

```
KillMode=process
```

```
Restart=always
```

```
RestartSec=5
```

```
# Having non-zero Limit*s causes performance problems due to accounting overhead
```

```
# in the kernel. We recommend using cgroups to do container-local accounting.
```

```
LimitNPROC=infinity
```

```
LimitCORE=infinity
```

```
LimitNOFILE=infinity
```

```
# Comment TasksMax if your systemd version does not supports it.
```

```
# Only systemd 226 and above support this version.
```

```
TasksMax=infinity
```

```
OOMScoreAdjust=-999
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
:wq!
```

```
# systemctl daemon-reload
```

```
# systemctl enable --now containerd
```

2- Installing runc

<https://github.com/opencontainers/runc/releases>

```
# wget https://github.com/opencontainers/runc/releases/download/v1.1.4/runc.amd64
```

```
# ls
```

```
containerd-1.6.8-linux-amd64.tar.gz runc.amd64
```

```
# install -m 755 runc.amd64 /usr/local/sbin/runc
```

3- Installing CNI plugins

<https://github.com/containernetworking/plugins/releases>

```
# wget https://github.com/containernetworking/plugins/releases/download/v1.1.1/cni-plugins-linux-amd64-v1.1.1.tgz
```

```
# ls
```

```
cni-plugins-linux-amd64-v1.1.1.tgz containerd-1.6.8-linux-amd64.tar.gz runc.amd64
```

```
# mkdir -p /opt/cni/bin
```

```
# tar Cxvf /opt/cni/bin cni-plugins-linux-amd64-v1.1.1.tgz
```

```
# sudo systemctl restart containerd
```

```
# systemctl status containerd
```

Kubernetes,

CRI-O

<https://github.com/cri-o/cri-o/blob/main/install.md#readme>

```
Operating system $OS
Centos 8 CentOS_8
CRI-O $VERSION
CRI-Ov1.23 1.23
```

```
# curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable.repo https://download.opensuse.org/repositories/devel:kubic:/libcontainers:/stable/CentOS_8/devel:kubic:libcontainers:stable.repo
# curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable:cri-o:1.23.repo https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable:cri-o:1.23/CentOS_8/devel:kubic:libcontainers:stable:cri-o:1.23.repo
```

```
# yum install cri-o -y
# yum install containernetworking-plugins -> it will install in CRI-O installation moment
# systemctl daemon-reload
# systemctl enable crio
# systemctl start crio
# systemctl status crio
```

Docker Engine

- 1- On each of your nodes, install **Docker** for your Linux distribution

to install Docker on RHEL **don't follow** <https://docs.docker.com/engine/install/#server>

to install Docker on RHEL **Follow** <https://faun.pub/how-to-install-simultaneously-docker-and-podman-on-rhel-8-centos-8-cb67412f321e>

```
# dnf config-manager --add-repo=https://download.docker.com/linux/centos/docker-ce.repo
# dnf config-manager --set-disabled docker-ce-stable
# rpm --install --nodeps --replacefiles --excludepath=/usr/bin/runc https://download.docker.com/linux/centos/8/x86_64/stable/Packages/containerd.io-1.4.9-3.1.el8.x86_64.rpm
# dnf install --enablerepo=docker-ce-stable docker-ce -y
# sudo systemctl enable --now docker
or
# systemctl enable docker.service
# systemctl start docker.service
# systemctl status docker.service
```

- 2- Install **cri-dockerd**, following the instructions in that source code repository

<https://www.mirantis.com/blog/how-to-install-cri-dockerd-and-migrate-nodes-from-dockershim/>
<https://github.com/Mirantis/cri-dockerd/releases>

```
# wget https://github.com/Mirantis/cri-dockerd/releases/download/v0.2.6/cri-dockerd-0.2.6.amd64.tgz
# tar xvf cri-dockerd-0.2.6.amd64.tgz
# sudo mv ./cri-dockerd /usr/local/bin/
# wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-docker.service
# wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-docker.socket
# mv cri-docker.socket cri-docker.service /etc/systemd/system/
# sed -i -e 's,/usr/bin/cri-dockerd,/usr/local/bin/cri-dockerd,' /etc/systemd/system/cri-docker.service
# systemctl daemon-reload
# systemctl enable cri-docker.service
# systemctl enable --now cri-docker.socket
# systemctl status cri-docker.socket
```

Installing kubeadm, kubelet and kubectl

kubeadm: the command to bootstrap the cluster.
kubelet: the component that runs on all of the machines in your cluster and does things like starting pods and containers.
kubectl: the command line util to talk to your cluster.

```
# cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-$basearch
enabled=1
gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
exclude=kubelet kubeadm kubectl
EOF
# sudo setenforce 0
# sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config

# sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes

# sudo systemctl enable --now kubelet
```

Kubernetes,

Creating a cluster with kubeadm

Kubernetes sensitive on SWAP space, it should be 0

```
# free -m
# swapoff -a
# free -m
```

on MasterNode

```
# hostname -i
192.168.0.191
```

Note: If kubelet is not running then user below four commands

1. sudo -i
2. swapoff -a
3. exit
4. strace -eopenat kubectl version

```
# # kubeadm init --apiserver-advertise-address=192.168.0.191 --cri-socket=/var/run/crio/crio.sock --v=5
```

```
.
.
.
```

Your Kubernetes control-plane has initialized **successfully!**

```
# mkdir -p $HOME/.kube
# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
# sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.0.191:6443 --token 7qu88a.4eqvd23ahd8yauus --discovery-token-ca-cert-hash sha256:6cc0c1339fbb546f658d23d091669483acc461533cca1befe17f9673c15e3163
```

on WorkerNodes

```
# kubeadm join 192.168.0.191:6443 --token 7qu88a.4eqvd23ahd8yauus --discovery-token-ca-cert-hash sha256:6cc0c1339fbb546f658d23d091669483acc461533cca1befe17f9673c15e3163 --cri-socket=/var/run/crio/crio.sock --v=5
```

NOTE: token doesn't work more than 23Hrs

on MasterNode

```
# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8smaster1.example.com	Ready	control-plane	5m32s	v1.25.3
k8sworker1.example.com	Ready	<none>	102s	v1.25.3
k8sworker3.example.com	Ready	<none>	112s	v1.25.3

```
# kubeadm completion bash >> ~/.bashrc
```

```
# kubectl completion bash >> ~/.bashrc
```

```
# echo "autocmd FileType yaml setlocal ai ts=2 sw=2 et cursorcolumn" >> ~/.vimrc
```