

18Oct2022

Day11

Kubernetes Resource Manager II

Policies-Resource Quotas

Policies-Limit Ranges

Monitoring resources through Metrix-Server

<https://kubernetes.io/docs/tasks/debug/debug-cluster/resource-metrics-pipeline/#metrics-server>

<https://github.com/kubernetes-sigs/metrics-server>

Metrics Server is a scalable, efficient source of container resource metrics for Kubernetes built-in autoscaling pipelines.

```
# kubectl top nodes
```

```
error: Metrics API not available
```

```
# kubectl top pod
```

```
error: Metrics API not available
```

Installation

<https://github.com/kubernetes-sigs/metrics-server>

```
# kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

```
# kubectl get deployments.apps -A
```

```
# kubectl get deployments.apps --namespace kube-system
```

```
# kubectl top nodes
```

```
Error from server (ServiceUnavailable): the server is currently unable to handle the request (get nodes.metrics.k8s.io)
```

```
# kubectl get deployments.apps --namespace kube-system
```

```
# kubectl get pods --namespace kube-system
```

```
# kubectl logs pods/metrics-server-678f4bf65b-jnm2f --namespace kube-system
```

to disable tls certification request,

```
# kubectl edit deployments.apps --namespace kube-system metrics-server
```

```
42 - --kubelet-preferred-address-types=InternalIP
```

->modify to 'InternalIP' remove others

```
43 - --kubelet-insecure-tls
```

->append this line

Note: If metric server doesn't start after above steps then also do this below line

Line below dnsPolicy: ClusterFirst

```
90 hostNetwork: true
```

```
:wq!
```

```
# kubectl top nodes
```

```
Error from server (ServiceUnavailable): the server is currently unable to handle the request (get nodes.metrics.k8s.io)
```

```
# kubectl rollout restart deployment --namespace kube-system metrics-server
```

```
# kubectl top nodes
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
master1.example.com	152m	7%	1684Mi	43%
node1.example.com	60m	6%	543Mi	28%

Resource Management for Pods and Containers

<https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/>

-When you specify a Pod, you can optionally specify how much of each resource a container need.

-The most common resources to specify are:

- 1- CPU
- 2- memory (RAM)

Resource types

CPU and memory are each a resource type

A resource type has a base unit.

CPU represents compute processing and is specified in units of Kubernetes CPUs.

CPU measurement is in millicores. with this measure format each **1CPU is splinted into 1000units(millicores)**

Memory is specified in units of **bytes**.

Memory can express as plain integer in B, K, M, G, T, P

Requests and Limits

Requests -> **minimum** requirement to start

when you specify the resource request for containers in a Pod, the kube-scheduler uses this information to decide which node to place the Pod on.

Limits -> **maximum** access to resources

when you specify a resource limit for a container, the kubelet enforces those limits so that the running container is not allowed to use more of that resource than the limit you set.

ex

```
# kubectl run nginx --image nginx -o yaml --dry-run=client >nginxpod.yaml
```

```
# vim nginxpod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
name: nginx
```

```
spec:
```

```
containers:
```

```
- image: nginx
```

```
name: nginx
```

```
resources:
```

```
requests:
```

```
cpu: "4m"
```

```
memory: "2Mi"
```

->to start, minimum

```
limits:
```

```
cpu: "8m"
```

```
memory: "4Mi"
```

->to make limit, maximum

```
:wq!
```

```
# kubectl top pods
NAME          CPU(cores)  MEMORY(bytes)
redisj-27764091-tfq6k  2m          2Mi
```

Policies

<https://kubernetes.io/docs/concepts/policy/>

Resource Quotas

<https://kubernetes.io/docs/concepts/policy/resource-quotas/>

-When several users or teams share a cluster with a fixed number of nodes, there is a concern that one team could use more than its fair share of resources.

-Resource quotas are a tool for administrators to address this concern.

-A resource quota, defined by a ResourceQuota object, provides constraints that limit aggregate resource consumption per **namespace**.

Implement Resource Quotas

- 1- create namespace
- 2- assign namespace to specific node/nodes
- 3- Enable Resource Quotas and define it on Namespace

1- create namespace

```
# kubectl create namespace testspace -o yaml --dry-run=client >testspace.yaml
```

```
# vim testspace.yaml
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: testspace
```

```
:wq!
```

```
# kubectl apply -f testspace.yaml
```

```
# kubectl get ns
```

```
testspace      Active 6s
```

2- assign namespace to specific node/nodes

2-1-enable PodNodeSelector

```
# vim /etc/kubernetes/manifests/kube-apiserver.yaml
```

```
20 - --enable-admission-plugins=NodeRestriction,PodNodeSelector
```

```
:wq!
```

```
# watch kubectl get nodes
```

2-2-label target node/nodes and namespace

```
# kubectl label nodes node2.example.com lbl=testspace
```

```
# kubectl get nodes node2.example.com --show-labels
```

```
# cat testspace.yaml
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: testspace
```

```
  annotations:
```

```
    scheduler.alpha.kubernetes.io/node-selector: lbl=testspace
```

```
:wq!
```

```
# kubectl apply -f testspace.yaml
```

verify

```
# kubectl run nginx --image nginx --namespace testspace
```

```
# kubectl get pods --namespace testspace -o wide
```

```
nginx 0/1 ContainerCreating 0 16s <none> node2.example.com
```

```
# kubectl delete pod --namespace testspace nginx --force --grace-period=0
```

3- define Resource Quotas on Namespace

```
# kubectl api-resources | grep -i "resource"
```

```
resourcequotas      quota      v1          true      ResourceQuota
```

```
# vim /etc/kubernetes/manifests/kube-apiserver.yaml
```

```
20 - --enable-admission-plugins=NodeRestriction,PodNodeSelector,ResourceQuota
```

```
:wq!
```

```
# watch kubectl get nodes
```

```
# cat testspacerq.yaml
```

```
kind: ResourceQuota
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: testspacerq
```

```
  namespace: testspace
```

```
spec:
```

```
  hard:
```

```
    pods: 5 requests.cpu:
```

```
      "10m"
```

```
    requests.memory: "8Mi"
```

```
    limits.cpu: "18m"
```

```
    limits.memory: "12Mi"
```

```
:wq!
```

```
# kubectl apply -f testspacerq.yaml --dry-run=client
```

```
# kubectl apply -f testspacerq.yaml
```

```
# kubectl get resourcequotas --namespace testspace
```

```
NAME      AGE  REQUEST      LIMIT
```

```
testspacerq 11s  pods: 0/5, requests.cpu: 0/10m, requests.memory: 0/8Mi  limits.cpu: 0/18m, limits.memory: 0/12Mi
```

```
# kubectl describe resourcequotas --namespace testspace
```

Verify

```
# kubectl create deployment dpl1 --image redis --replicas 6 --namespace testspace
# kubectl run redis --image redis --namespace testspace -o yaml --dry-run=client >redispod.yaml
```

```
# vim nginxpod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
    name: nginxpod
    namespace: testspace
spec:
  containers:
    - image: nginx
      name: nginxcnt
      resources:
        requests:
          cpu: "11m"
          memory: "6Mi"
        limits:
          cpu: "19m"
          memory: "8Mi"
```

```
:wq!
```

```
# kubectl create -f redispod.yaml
```

Error from server (Forbidden): error when creating "redispod.yaml": pods "redispod" is forbidden: exceeded quota: testspacerq, requested: limits.cpu=19m,requests.cpu=11m, used: limits.cpu=0,requests.cpu=0, limited: limits.cpu=18m,requests.cpu=10m

```
# cat nginxpod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
    name: nginxpod
    namespace: testspace
spec:
  containers:
    - image: nginx
      name: nginxcnt
      resources:
        requests:
          cpu: "3m"
          memory: "4Mi"
        limits:
          cpu: "6m"
          memory: "12Mi"
```

```
:wq!
```

```
# kubectl create -f redispod.yaml
```

```
# kubectl get pods --namespace testspace nginxpod
```

```
# kubectl describe pods --namespace testspace nginxpod
```

```
# kubectl describe resourcequotas --namespace testspace testspacerq
```

Name:	testspacerq	
Namespace:	testspace	
Resource	Used	Hard
limits.cpu	6m	18m
limits.memory	12Mi	12Mi
pods	1	5
requests.cpu	3m	10m
requests.memory	4Mi	8Mi

Limit Ranges

<https://kubernetes.io/docs/concepts/policy/limit-range/>

-By default, containers run with unbounded compute resources on a Kubernetes cluster.

-Using Kubernetes resource quotas, administrators can restrict consumption and creation of cluster resources (such as CPU, memory, and persistent storage) within a specified namespace.

issue

-Within a namespace, a Pod can consume as much CPU and memory as is allowed by the ResourceQuotas that apply to that namespace. -> **issue is here**

-As a cluster operator, or as a namespace-level administrator, you might also be concerned about making sure that a single object cannot monopolize all available resources within a namespace.

solution

A **LimitRange** is a policy to constrain the resource allocations (limits and requests) that you can specify for each applicable object kind (such as Pod or PersistentVolumeClaim) in a namespace.

Enable **LimitRange**

by default, is Enabled

Implementing Limit Ranges

- 1- define Metrics Server
- 2- assign namespace to specific node/nodes
- 3- implement Resource Quotas
- 4- implement LimitRanges

```
# kubectl top nodes
```

```
# cat testspace.yaml
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: testspace
```

```
  annotations:
```

```
    scheduler.alpha.kubernetes.io/node-selector: lbl=testspace
```

```
# kubectl get nodes node2.example.com --show-labels
```

```
node2.example.com Ready <none> 82d v1.24.3 beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=node2.example.com,kubernetes.io/os=linux,lbl=testspace
```

```
# kubectl get resourcequotas --namespace testspace testspacerq
```

```
NAME      AGE  REQUEST      LIMIT
```

```
testspacerq 41m  pods: 0/5, requests.cpu: 0/10m, requests.memory: 0/8Mi limits.cpu: 0/18m, limits.memory: 0/12Mi
```

```
# kubectl api-resources | grep -i "limit"
```

```
limitranges      limits      v1          true      LimitRange
```

```
# vim testspacelr.yaml
```

```
kind: LimitRange
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: testspacelr
```

```
namespace: testspace
```

```
spec:
```

```
  limits:
```

```
    - type: Container      ->Object Name
```

```
      min:
```

```
        cpu: "3m"          ->min amount of CPU that single container can request it if defines it inside Pod.
```

```
        memory: "5Mi"      ->min amount of MEMORY that single container can request it if defines it inside Pod.
```

```
      max:
```

```
        cpu: "5m"          ->max amount of CPU that single container can achieve it if defines it inside Pod.
```

```
        memory: "14Mi"     ->max amount of MEMORY that single container can achieve it if defines it inside Pod.
```

```
      defaultRequest:      ->requests
```

```
        cpu: "4m"          ->min amount of CPU that single container can request it if doesn't define it inside Pod.
```

```
        memory: "6Mi"      ->min amount of MEMORY that single container can request it if doesn't define it inside Pod.
```

```
      default:            ->limits
```

```
        cpu: "4m"          ->max amount of CPU that single container can achieve it if doesn't define it inside Pod.
```

```
        memory: "12Mi"     ->max amount of MEMORY that single container can achieve it if doesn't define it inside Pod.
```

```
:wq!
```

```
# kubectl apply -f testspacelr.yaml --dry-run=client
```

```
# kubectl apply -f testspacelr.yaml
```

```
# kubectl describe limitranges --namespace testspace
```

```
# kubectl describe resourcequotas --namespace testspace
```

verify

```
# cat nginxpod2.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  labels:
```

```
    run: nginx
```

```
  name: nginxpod
```

```
  namespace: testspace
```

```
spec:
```

```
  containers:
```

```
    - image: nginx
```

```
      name: nginxcnt
```

shouldn't less than min-cpu

shouldn't less than min-memory

shouldn't greater than min-cpu

shouldn't greater than min-memory

```
# cat nginxpod1.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginxpod
  namespace: testspace
spec:
  containers:
  - image: nginx
    name: nginxcnt
    resources:
      requests:
        cpu: "3m"
        memory: "5Mi"
      limits:
        cpu: "5m"
        memory: "12Mi"
```