

20Oct2022

Day12

Kubernetes Network

Kubernetes Service

Pods

Pods are the smallest deployable units of computing that you can create and manage in Kubernetes. Pods natively provide two kinds of shared resources for their constituent containers: **networking** and **storage**.

after create K8s cluster, by default one **network plug-in** be available and has been installed, name is CNI
when create POD, ip will set on POD not container, not deployment, not ReplicaSet,...



what is Container Network Interface-CNI

<https://github.com/containernetworking/cni>

CNI-Container Network Interface, a **CNCF** project, consists of a specification and libraries for writing plugins to configure network interfaces in Linux containers, along with a number of supported plugins.

ip a s

```
4: cni0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    inet 10.85.0.1/16 brd 10.85.255.255 scope global cni0
```

cat /etc/cni/net.d/100-crio-bridge.conf

CNCF-CNI advantage

-default network plugin

CNCF-CNI disadvantage

-application accessible from hosted node, not from other nodes in cluster

```
# kubectl run pod1 --image nginx -o yaml --dry-run=client > pod1.yaml
```

```
# vim pod1.yaml
```

```
# cat pod1.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  labels:
```

```
    run: pod1
```

```
  name: pod1
```

```
spec:
```

```
  containers:
```

```
    - image: nginx
```

```
      name: nginxcnt
```

```
# kubectl apply -f pod1.yaml
```

```
# kubectl get pods -o wide
```

```
pod1 1/1 Running 0 16m 10.85.0.2 node2.example.com
```

```
root@master:~# curl 10.85.0.2
```

```
curl: (7) Failed to connect to 10.85.0.2 port 80: No route to host
```

```
root@master:~# ssh root@node2.example.com
```

```
root@node2:~# curl 10.85.0.2
```

```
<h1>Welcome to nginx!</h1>
```

```
root@node2:~# logout
```

```
root@master:~#
```

Installing Addons

<https://kubernetes.io/docs/concepts/cluster-administration/addons/#networking-and-network-policy>

<https://github.com/containernetworking/cni>

Calico

<https://github.com/projectcalico/calico>

<https://projectcalico.docs.tigera.io/about/about-calico>

its opensource CNI plugin for Kubernetes.

Calico is a widely adopted, battle-tested open-source networking and network security solution for Kubernetes, virtual machines, and bare-metal workloads.

Calico provides two major services for Cloud Native applications:

1-Network connectivity between workloads.

2-Network security policy enforcement between workloads.

Implement Calico on K8s Cluster

<https://docs.projectcalico.org/manifests/calico.yaml>

```
# wget https://docs.projectcalico.org/manifests/calico.yaml
```

```
# ls
```

```
calico.yaml
```

```
# kubectl apply -f calico.yaml
```

```
# watch kubectl get pods -A -A ->whole namespace
```

```
# watch kubectl get pods --namespace kube-system
```

```
# kubectl get pods -o wide
```

```
# kubectl delete pod pod1 --force --grace-period=0
```

```
# kubectl create -f pod1.yaml
```


```
# kubectl get pods -o wide
```

```
pod1 1/1 Running 0 7m45s 172.16.221.1 node2.example.com
```

NOW, from every node in cluster container is accessible

```
root@node3:~# curl 172.16.221.1
```

```
<h1>Welcome to nginx!</h1>
```



issue will happen when create deployment

```
# kubectl create deployment dpl1 --image redis --replicas=3 -o yaml --dry-run=client >dpl1.yaml
# kubectl create -f dpl1.yaml
# kubectl get deployments.apps -o wide
# kubectl get pods -o wide
dpl1-7b5786d846-6cqr6 1/1 Running 0 24s 172.16.221.2 node2.example.com
dpl1-7b5786d846-bztd8 1/1 Running 0 24s 172.16.11.65 node1.example.com
dpl1-7b5786d846-dxsp8 1/1 Running 0 24s 172.16.178.193 node3.example.com
```

Points

-after create deployment, it depends on pods inside it, ip will assign to pod

-next issue is when a pod inside deployment will delete again new one will generate immediately but with new ip from CNI. IP, its temporary.

above issues will solve by SERVICE in K8s cluster

Service

<https://kubernetes.io/docs/concepts/services-networking/service/>

An abstract way to expose an application running on a set of Pods as a network service.

```
# kubectl create service --help
```

Create a service using a specified subcommand.

Available Commands:

clusterip	Create a ClusterIP service	->exposes the service just inside the cluster
externalname	Create an ExternalName service	
loadbalancer	Create a LoadBalancer service	
nodeport	Create a NodePort service	->application accessible from outside through static, permanent ip

-create NGINX deployment first with 3replicas

```
# kubectl get deployment
# kubectl get pods
# kubectl expose deployment dpl1 --name dpl1int --port 80 --protocol TCP --type ClusterIP -o yaml --dry-run=client >dpl1int.yaml
# cat dpl1int.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: dpl1
  name: dpl1int
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: dpl1
  type: ClusterIP
status:
  loadBalancer: {}
# kubectl expose deployment dpl1 --name dpl1int --port 80 --protocol TCP --type ClusterIP
# kubectl get svc
dpl1int ClusterIP 10.110.147.73 <none> 80/TCP 32s
# kubectl describe svc dpl1int
# kubectl expose deployment dpl1 --name dpl1ext --port 80 --protocol TCP --type NodePort -o yaml --dry-run=client >dpl1ext.yaml
# cat dpl1ext.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: dpl1
  name: dpl1ext
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: dpl1
  type: NodePort
status:
  loadBalancer: {}
# kubectl apply -f dpl1ext.yaml
# kubectl get svc
dpl1ext NodePort 10.97.239.101 <none> 80:30991/TCP 4s
dpl1int ClusterIP 10.110.147.73 <none> 80/TCP 6m8s
```

NOTE: NodePort type its between **30000 to 32767**

Verify

open web browser and type

<http://192.168.0.240:30991/>

Welcome to nginx!