

21,25Oct2022


Day13, 14

Kubernetes Storage

Pods

Pods are the smallest deployable units of computing that you can create and manage in Kubernetes. Pods natively provide two kinds of shared resources for their constituent containers: **Networking** and **Storage**.

after create K8s cluster, by default one **network plug-in** be available and has been installed, name is **CNCF-CNI PLUGIN** when create POD, Ip will set on it not container/s



Networking

what is Container Network Interface-CNI

<https://github.com/containernetworking/cni>



CNI-Container Network Interface, a **CNCF** project, consists of a specification and libraries for writing plugins to configure network interfaces in Linux containers, along with a number of supported plugins.

Storage

Container Storage Interface-CSI

container storage is ephemeral.

needs map container's storage to outside of pod.

ex

```
$ kubectl run pod1 --image=nginx
$ kubectl get pods
$ kubectl exec -it pod1 -- ls /
$ kubectl exec -it pod1 -- cat /usr/share/nginx/html/index.html
$ kubectl exec -it pod1 -- /bin/bash
root@pod1:/# curl localhost
Welcome to Nginx!
root@pod1:/# exit
$ kubectl delete pod pod1 --force --grace-period=0
```

Now, Pod has been deleted, data gone. to solve this issue, need practice on Storage in K8s

Storage

<https://kubernetes.io/docs/concepts/storage/>

in Kubernetes Pod will get connection to external storage through **PV** and **PVC**

Persistent Volume-PV

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

-The PersistentVolume subsystem provides an API for users and administrators that abstracts details of how storage is provided from how it is consumed.

-A PersistentVolume-PV is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes.

-It is a resource in the cluster just like a node is a cluster resource.

-its cluster-based

Persistent Volume Claim-PVC

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

A PersistentVolumeClaim-PVC is a request for storage by a user.

It is similar to a Pod;

Pods consume node resources and PVCs consume PV resources.

Pods can request specific levels of resources (CPU and Memory) and Claims can request specific size and access modes from PV

-namespace/project-based resource

Implement Storage in K8s Cluster

step1-create PV
 step2-create PVC
 step3-maps PVC to PV
 step4-maps container/s to PV through PVC

step1-create PV

step2-create PVC

step3-maps PVC to PV

-Reclaiming

when a user is done with their volume, they can delete the PVC objects from the API that allows reclamation of the resource. The reclaim policy for a PersistentVolume tells the cluster what to do with the volume after it has been released of its claim. Currently, volumes can either be Retained, Recycled, or Deleted.

1-Retain

The Retain reclaim policy allows for manual reclamation of the resource.

When the PersistentVolumeClaim-PVC is deleted, the PersistentVolume-PV still exists and the volume is considered "released".

2-Delete

deletion removes both the PersistentVolume-PV object from Kubernetes, as well as the associated storage asset in the external infrastructure, such as an AWS EBS, GCE PD, Azure Disk.

3-Recycle

If supported by the underlying volume plugin, the Recycle reclaim policy performs a basic scrub (rm -rf /thevolume/*) on the volume and makes it available again for a new claim.

-Access Modes

A PersistentVolume can be mounted on a host in any way supported by the resource provider.

1-ReadWriteOnce-RWO

the volume can be mounted as read-write by a single node.

ReadWriteOnce access mode still can allow multiple pods to access the volume when the pods are running on the same node.

2-ReadWriteMany-RWX

the volume can be mounted as read-write by many nodes.

3-ReadOnlyMany-ROX

the volume can be mounted as read-only by many nodes.

4-ReadWriteOncePod-RWOP

the volume can be mounted as read-write by a single Pod.

Use ReadWriteOncePod access mode if you want to ensure that only one pod across whole cluster can read that PVC or write to it.

```
$ kubectl api-resources | grep -i "pv"
```

```
persistentvolumes    pv    v1    false    PersistentVolume
```

```
# vim pv1.yaml
```

```
kind: PersistentVolume
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: pv1
```

```
spec:
```

```
  capacity:
```

```
    storage: 1Gi
```

```
  accessModes:
```

```
    - ReadWriteMany
```

```
  persistentVolumeReclaimPolicy: Retain
```

```
  hostPath:
```

```
    path: /mnt/disk1
```

```
:wq!
```

```
# kubectl apply -f pv1.yaml --dry-run=client
```

```
# kubectl apply -f pv1.yaml
```

```
# kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
pv1	1Gi	RWX	Retain	Available				2m47s

```
# vim pvc1.yaml
```

```
kind: PersistentVolumeClaim
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: pvc1
```

```
spec:
```

```
  resources:
```

```
    requests:
```

```
      storage: 2Gi
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

```
:wq!
```

```
# kubectl create -f pvc1.yaml
```

```
# kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
pvc1	Pending					4s

NOTE: PV, Cluster-based resource
 PVC, namespace-based resource

NOTE

to bind PVC to PV first permission will check on PV's then size.

if Permission passes, size should be acceptable, example PVC with higher size than PV wont bind. **check above example**

Now,

```
# kubectl edit pv pv1
```

```
22 storage: 3Gi
```

```
:wq!
```

PV size has been changed/updated

```
# kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
pvc1	Bound	pv1	3Gi	RWX		58s

```
# kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
pv1	3Gi	RWX	Retain	Bound	default/pvc1			80s

ex of change in size and permission

```
# vim pv1.yaml
```

```
kind: PersistentVolume
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: pv1
```

```
spec:
```

```
  capacity:
```

```
    storage: 1Gi
```

```
  accessModes:
```

```
    - ReadWriteMany
```

```
  persistentVolumeReclaimPolicy: Retain
```

```
  hostPath:
```

```
    path: /mnt/disk1
```

```
:wq!
```

```
# vim pvc1.yaml
```

```
kind: PersistentVolumeClaim
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: pvc1
```

```
spec:
```

```
  resources:
```

```
    requests:
```

```
      storage: 2Gi
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

```
:wq!
```

```
# kubectl apply -f pv1.yaml
```

```
# kubectl apply -f pvc1.yaml
```

```
# kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
pv1	3Gi	RWX	Retain	Released				20m

```
# kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
pvc1	Pending					8m21s

step4-maps container/s to PV through PVC

```
# kubectl run pod1 --image=nginx -o yaml --dry-run=client >pod1.yaml
```

```
# vim pod1.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  labels:
```

```
    run: pod1
```

```
  name: pod1
```

```
spec:
```

```
  containers:
```

```
    - image: nginx
```

```
      name: nginxcnt
```

```
  volumeMounts:
```

```
    - name: nginxstr
```

```
      mountPath: "/usr/share/nginx/html"
```

```
  volumes:
```

```
    - name: nginxstr
```

```
      persistentVolumeClaim:
```

```
        claimName: pvc1
```

```
:wq!
```

NOTE: container through volumeMount's name(**nginxstr**) will connect to PVC volumes name. both name should be same.

```
# kubectl apply -f pod1.yaml --dry-run=client
```

```
# kubectl get pv
```

```
# kubectl get pvc
```

```
# kubectl get po
```

```
# kubectl expose pod pod1 --name=podint --port=80 --protocol=TCP --type=ClusterIP
# kubectl get svc
# kubectl get po -o wide
NAME READY STATUS RESTARTS AGE IP NODE
pod1 1/1 Running 0 7m13s 172.18.184.3 k8sworker1.example.com
# curl 10.98.40.149
<html>
<head><title>403 Forbidden</title></head>
<body>
<center><h1>403 Forbidden</h1></center>
# ssh root@k8sworker1.example.com ls /mnt
root@k8sworker1.example.com's password:
disk1
# ssh root@k8sworker1.example.com
root@k8sworker1.example.com's password: ****
Last login: Tue Oct 25 19:44:38 2022 from 10.0.0.12
[root@k8sworker1 ~]# echo "Hello!" >/mnt/disk1/index.html
[root@k8sworker1 ~]# exit
# curl 10.98.40.149
Hello!
```

Now, delete POD

```
# kubectl delete pod pod1 --force
# curl 10.98.40.149
curl: (7) Failed to connect to 10.98.40.149 port 80: Connection refused
# ssh root@k8sworker1.example.com ls /mnt/disk1
root@k8sworker1.example.com's password: ****
index.html
[root@k8smaster1 ~]# ssh root@k8sworker1.example.com cat /mnt/disk1/index.html
root@k8sworker1.example.com's password: ****
Hello!
```

Now, delete the PVC

```
# kubectl get pv pv1 -o yaml | grep -i "persistentvolumereclaimpolicy"
  persistentVolumeReclaimPolicy: Retain
# kubectl delete pvc pvc1
# ssh root@k8sworker1.example.com ls /mnt/disk1
root@k8sworker1.example.com's password: ****
index.html
[root@k8smaster1 ~]# ssh root@k8sworker1.example.com cat /mnt/disk1/index.html
root@k8sworker1.example.com's password: ****
Hello!
```