

26Oct2022

Day15

## Kubernetes Storage II-access to remote storage

### Storage

<https://kubernetes.io/docs/concepts/storage/>

in Kubernetes Pod will get connection to external storage through **PV** and **PVC**

#### Persistent Volume-PV

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

-The PersistentVolume subsystem provides an API for users and administrators that abstracts details of how storage is provided from how it is consumed.

-A PersistentVolume-PV is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes.

-It is a resource in the cluster just like a node is a cluster resource.

-its cluster-based

#### Persistent Volume Claim-PVC

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>


A PersistentVolumeClaim-PVC is a request for storage by a user.

It is similar to a Pod:

-Pods consume node resources and PVCs consume PV resources.

-Pods can request specific levels of resources (CPU and Memory) and Claims can request specific size and access modes from PV

-namespace/project-based resource



### Implement Storage in K8s Cluster

step1-create PV

step2-create PVC

step3-maps PVC to PV

step4-maps container/s to PV through PVC

### PV, PVC parameters explanation

#### -Reclaiming

when a user is done with their volume, they can delete the PVC objects from the API that allows reclamation of the resource.

The reclaim policy for a PersistentVolume tells the cluster what to do with the volume after it has been released of its claim.

Currently, volumes can either be Retained, Recycled, or Deleted.

#### 1-Retain

The Retain reclaim policy allows for manual reclamation of the resource.

When the PersistentVolumeClaim-PVC is deleted, the PersistentVolume-PV still exists and the volume is considered "released".

#### 2-Delete

deletion removes both the PersistentVolume-PV object from Kubernetes, as well as the associated storage asset in the external infrastructure, such as an AWS EBS, GCE PD, Azure Disk.

#### 3-Recycle

If supported by the underlying volume plugin, the Recycle reclaim policy performs a basic scrub (`rm -rf /thevolume/*`) on the volume and makes it available again for a new claim.

#### -Access Modes

A PersistentVolume can be mounted on a host in any way supported by the resource provider.

#### 1-ReadWriteOnce-RWO

the volume can be mounted as read-write by a single node.

ReadWriteOnce access mode still can allow multiple pods to access the volume when the pods are running on the same node.

#### 2-ReadWriteMany-RWX

the volume can be mounted as read-write by many nodes.

#### 3-ReadOnlyMany-ROX

the volume can be mounted as read-only by many nodes.

#### 4-ReadWriteOncePod-RWOP

the volume can be mounted as read-write by a single Pod.

Use ReadWriteOncePod access mode if you want to ensure that only one pod across whole cluster can read that PVC or write to it.

## Attach Remote Storage to K8s Deployment

Remote Storages:

### 1-File-Based Storage Protocol

1-1-NFS

1-2-CIFS

### 2-Block-Based Storage Protocol

2-1-iSCSI/Tragetcli

### 1-File-Based Storage Protocol

1-1-NFS Network File Sharing/System

### Implement NFS on RHEL

Install NFS on RHEL Linux machine out of K8s Cluster but accessible by K8s Cluster.

login to NFS-Server Host

u: root

p: redhat

#### -Basic Configuration

##### -Hostname

```
# hostnamectl set-hostname generic.example.com
```

```
# hostnamectl
```

##### -Network

```
# nmcli connection show
```

```
enp0s3 5abdbad1-f4d4-48be-bc18-2c552b4aea9a ethernet enp0s3
```

```
virbr0 5abdbad1-f4d4-48be-bc18-2c552b4aea9a ethernet virbr0
```

**NOTE: don't delete virbr0**

```
# nmcli connection delete "enp0s3"
```

```
# nmcli connection reload
```

```
# ifconfig
```

```
# nmcli connection add con-name "enp0s3" type ethernet autoconnect yes ifname enp0s3 ipv4.addresses "192.168.29.110/24" ipv4.dns "192.168.29.1"
```

```
ipv4.gateway "192.168.29.1" +ipv4.dns "8.8.8.8" +ipv4.dns "8.8.4.4" ipv4.method manual
```

```
# nmcli connection reload
```

```
# reboot
```

```
# nmcli connection show
```

```
# ifconfig
```

```
# ping 8.8.8.8
```

##### -Storage

```
# lsblk
```

```
sdb      8:16 0 5G 0 disk
```

```
# pvcreate /dev/sdb
```

```
# vgcreate vg1 /dev/sdb
```

```
# lvcreate -n lv1 -l 100%FREE vg1
```

```
# pvs
```

```
# vgs
```

```
# lvs
```

```
# mkfs.xfs /dev/mapper/vg1-lv1
```

```
# mkdir /mnt/disk1
```

```
# ls -ld /mnt/disk1/
```

```
drwxr-xr-x 2 root root 6 Oct 26 19:52 /mnt/disk1/
```

```
# chmod 757 /mnt/disk1/
```

```
# ls -ld /mnt/disk1/
```

```
drwxr-xrwx 2 root root 6 Oct 26 19:52 /mnt/disk1/
```

```
# blkid
```

```
/dev/mapper/vg1-lv1: UUID="c766c56d-f5a5-45e9-8ffa-443d931eaa05" BLOCK_SIZE="512" TYPE="xfs"
```

```
# echo "/dev/mapper/vg1-lv1 /mnt/disk1 xfs defaults 0 0" >>/etc/fstab
```

```
# mount -a
```

```
# df -hT
```

```
/dev/mapper/vg1-lv1 xfs 5.0G 68M 5.0G 2% /mnt/disk1
```

## -Config NFS-server

install NFS-Server packages

### Local Repository

```
# lsblk
sr0      11:0  1 10.7G 0 rom
# mkdir /media/cdrom
# echo "/dev/sr0 /media/cdrom iso9660 defaults 0 0" >>/etc/fstab
# mount -a
# df -hT
/dev/mapper/vg1-lv1 xfs      5.0G 68M 5.0G 2% /mnt/disk1
/dev/sr0            iso9660 11G  11G  0   100% /media/cdrom
# vim /etc/yum.repos.d/redhat.repo
[App]
name=AppStream
baseurl=file:///media/cdrom/AppStream/
gpgcheck=0
enabled=1
[Base]
name=BaseOS
baseurl=file:///media/cdrom/BaseOS/
gpgcheck=0
enabled=1
:wq!
# yum repolist
```

### EPEL repository

```
https://docs.fedoraproject.org/en-US/epel/
# dnf install https://dl.fedoraproject.org/pub/epel/epel-release-latest-8.noarch.rpm -y
# yum repolist
```

### RedHat subscription Manager

```
# subscription-manager register
username: naghval                                ->RHN account
password: ****
# subscription-manager attach --auto
# subscription-manager release --set=8.6          ->its option
# yum repolist
```

### Install NFS-Server

```
disable SELinux and firewall
# sestatus
# getenforce
Enforcing
# vim /etc/selinux/config
SELINUX=disabled
:wq!
# setenforce 0
# systemctl disable firewalld.service && systemctl stop firewalld.service && systemctl mask firewalld.service
# reboot
# sestatus
SELinux status:      disabled
# systemctl status firewalld.service
# yum list nfs-utils rpcbind
# yum install nfs-utils rpcbind -y
# systemctl enable nfs-server.service
# systemctl start nfs-server.service
# systemctl status nfs-server.service
# df -hT
/dev/mapper/vg1-lv1 xfs      5.0G 68M 5.0G 2% /mnt/disk1
# vim /etc/exports
<export storage> <to whom should be accessible>(<permissions>)
/mnt/disk1 *.example.com(rw,sync)
/mnt/disk1 192.168.29.0/24(rw,sync)
:wq!
# systemctl restart nfs-server.service
```

->its DNS example/sample

### verify

local(on NFS-Server only)

```
# exportfs -rva
exporting 192.168.29.0/24:/mnt/disk1
```

global(client-side)

```
# showmount -e 192.168.29.110
Export list for 192.168.29.110:
/mnt/disk1 192.168.29.0/24
```

## Install NFS-Client

config K8s Cluster Nodes as NFS-Client  
it depends on your K8s Cluster nodes OS

### RHEL

```
# yum install nfs-utils -y
# showmount -e 192.168.29.110
```

### UBUNTU

```
# apt install nfs-common -y
# ping 192.168.29.110 -c 2
# showmount -e 192.168.29.110
Export list for 192.168.29.110:
/mnt/disk1 192.168.29.0/24
```

**NOTE:** install 'nfs-common' on whole K8s Cluster Nodes.

## Continue on MasterNode

```
# kubectl api-resources | grep -i "pv"
# vim pv1.yaml
kind: PersistentVolume
apiVersion: v1
metadata:
  name: pv1
spec:
  capacity:
    storage: 3Gi
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Recycle
  nfs:
    path: /mnt/disk1
    server: 192.168.29.110
---
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: pvc1
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 2Gi
:wq!
# kubectl apply -f pv1.yaml --dry-run=client
# kubectl apply -f pv1.yaml
# kubectl get pv
pv1 3Gi RWX Recycle Bound default/pvc1 3s
# kubectl get pvc
pvc1 Bound pv1 3Gi RWX 7s
# kubectl create deployment dpl1 --image nginx --replicas=3 -o yaml --dry-run=client
# kubectl create deployment dpl1 --image nginx --replicas=3 -o yaml --dry-run=server
# kubectl create deployment dpl1 --image nginx --replicas=3 -o yaml --dry-run=client > dpl1.yaml
# vim dpl1.yaml
# vim dpl1.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: dpl1
  name: dpl1
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginxapp
  template:
    metadata:
      labels:
        app: nginxapp
    spec:
      containers:
        - image: nginx
          name: nginxcnt
          volumeMounts:
            - name: nginxstr
              mountPath: "/usr/share/nginx/html"
      volumes:
        - name: nginxstr
          persistentVolumeClaim:
            claimName: pvc1
:wq!
```

```
# kubectl apply -f dpl1.yaml --dry-run=client
# kubectl apply -f dpl1.yaml
# kubectl get deployments.apps
dpl1 3/3 3 3 2m13s
# kubectl get po
dpl1-85f74f8c75-h2f2z 1/1 Running 0 2m53s 172.16.221.1 node2.example.com
dpl1-85f74f8c75-mtccs 1/1 Running 0 2m53s 172.16.206.1 node4.example.com
dpl1-85f74f8c75-nsczg 1/1 Running 0 2m53s 172.16.11.65 node1.example.com
# curl 172.16.221.1
<html>
<head><title>403 Forbidden</title></head>
```

### ClusterIP

```
# kubectl expose deployment dpl1 --name dpl1int --port 80 --protocol TCP --type ClusterIP
```

### NodePort

```
# kubectl expose deployment dpl1 --name dpl1ext --port 80 --protocol TCP --type NodePort
# kubectl get svc
dpl1ext NodePort 10.105.142.163 <none> 80:32165/TCP 17s
dpl1int ClusterIP 10.106.59.137 <none> 80/TCP 52s
```

### verify

```
# curl 10.106.59.137
<html>
<head><title>403 Forbidden</title></head>
# hostname -i
192.168.29.104
open web browser
http://192.168.29.104:32165
403 Forbidden
now, back to NFS-Server and put customer NGINX server content in to DirectoryIndex file
root@master1:~# ssh root@192.168.29.110
password: redhat
[root@generic ~]# echo "Hellooooooooooooooooo!" >/mnt/disk1/index.html
# cat /mnt/disk1/index.html
Hellooooooooooooooooo!
[root@generic ~]# logout
root@master1:~# # curl 10.106.59.137
Hellooooooooooooooooo!
open web browser
http://192.168.29.104:32165
Hellooooooooooooooooo!
```