

27Oct2022

Day16

ETCD

Horizontal Pod Autoscaling-HPA

etcd, K8s Database

<https://etcd.io/>

what is ETCD?

-A distributed, reliable key-value store for the most critical data of a distributed system.

-etcd is a strongly consistent, distributed key-value store that provides a reliable way to store data that needs to be accessed by a distributed system or cluster of machines

on **MasterNode**

hostname

master1.example.com

hostname -i

192.168.29.104

ls /etc/kubernetes/manifests/

etcd.yaml kube-apiserver.yaml kube-controller-manager.yaml kube-scheduler.yaml

vim /etc/kubernetes/manifests/etcd.yaml

etcd port number -> 2379, 2380

etcd database location -> /var/lib/etcd

ls /var/lib/etcd/member/snap/

db

etcd pub/private key + ca-certification

ls /etc/kubernetes/pki/etcd/

ca.crt ca.key healthcheck-client.crt healthcheck-client.key peer.crt peer.key server.crt server.key

Interact with ETCD

<https://github.com/etcd-io/etcd/releases>

mkdir etcd

cd etcd

wget <https://github.com/etcd-io/etcd/releases/download/v3.5.5/etcd-v3.5.5-linux-amd64.tar.gz>

pwd

/root/etcd

ls

etcd-v3.5.5-linux-amd64.tar.gz

tar xfv etcd-v3.5.5-linux-amd64.tar.gz

cp -r etcd-v3.5.5-linux-amd64/etcd* /usr/local/bin/

ls /usr/local/bin/

etcdctl --help

etcd --version

etcd Version: 3.5.5

Git SHA: 19002cfc6

Go Version: go1.16.15

Go OS/Arch: linux/amd64

etcdctl version

etcdctl version: 3.5.5

API version: 3.5

Backup ETCD

step1-make activity on K8s cluster

kubectl create deployment dpl1 --image nginx --replicas 3

kubectl create deployment dpl2 --image redis --replicas 3

kubectl create namespace space1

kubectl get deployments.apps

kubectl get ns

kubectl get pods

step2-take backup

to take backup need pub-key, private-key and certification but, here we use K8s cluster etcd own pub-key, private-key and certification.

etcdctl --help

etcdctl version

API version: 3.5

ls /etc/kubernetes/pki/etcd/

ca.crt ca.key healthcheck-client.crt healthcheck-client.key peer.crt peer.key server.crt server.key

backup

ETCDCTL_API=3 etcdctl snapshot save --endpoints=127.0.0.1:2379 --cert="/etc/kubernetes/pki/etcd/server.crt" --key="/etc/kubernetes/pki/etcd/server.key" --cacert="/etc/kubernetes/pki/etcd/ca.crt" /opt/etcd27Oct22.db

ls /opt

etcd27Oct22.db

du -h /opt/etcd27Oct22.db

5.4M /opt/etcd27Oct22.db

status

```
# etcdctl snapshot status /opt/etcd27Oct22.db
# etcdctl snapshot status /opt/etcd27Oct22.db
b73d8ba7, 7570, 1810, 5.6 MB
# etcdctl snapshot status /opt/etcd27Oct22.db --write-out="table"
# etcdctl snapshot status /opt/etcd27Oct22.db --write-out="table"
```

```
+-----+-----+-----+
|HASH| |REVISION| |TOTAL KEYS| |TOTAL SIZE|
+-----+-----+-----+
| b73d8ba7 | 7570 | 1810 | 5.6 MB |
+-----+-----+-----+
```

Restore ETCD

step1-delete some resources

```
# kubectl get deployments.apps
# kubectl get ns
# kubectl delete deployments.apps dpl1 dpl2 --force --grace-period=0
# kubectl delete namespaces space1 --force --grace-period=0
```

restore

```
current ->/var/lib/etcd/member/snap/db
backup ->/opt/etcd27Oct22.db
```

verify

```
# cat /etc/kubernetes/manifests/etcd.yaml
- hostPath:
  path: /var/lib/etcd
```

->current etcd db location connecting to K8s cluster kube-api

```
# ETCDCTL_API=3 etcdctl snapshot restore --endpoints=127.0.0.1:2379 --cert="/etc/kubernetes/pki/etcd/server.crt" --
key="/etc/kubernetes/pki/etcd/server.key" --cacert="/etc/kubernetes/pki/etcd/ca.crt" --data-dir=/var/lib/newetcd/ /opt/etcd27Oct22.db
# ll /var/lib
drwx----- 3 root root 4096 Oct 27 13:39 etcd/
drwx----- 3 root root 4096 Oct 27 14:45 newetcd/
# vim /etc/kubernetes/manifests/etcd.yaml
76 - hostPath:
```

```
77 path: /var/lib/newetcd
:wq!
```

->update etcd about new db location

```
# kubectl get nodes
# kubectl get deployments.apps
# kubectl get ns
```


Horizontal Pod Autoscaling-HPA

<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>

- a **HorizontalPodAutoscaler** automatically updates a workload resource, with the aim of automatically scaling the workload to match demand.
- Horizontal scaling means that the response to increased load is to deploy more Pods.

Procedure

- config Metrics-Server on K8s cluster
- deploy an application, specify CPU and MEMORY on deployment.
- Now**, if Metrics-Server detects suppose 90% of CPU used then new pod will create.



Implement HPA on K8s Cluster

CPU-based

steps

- enable Metrics-Server
- deploy an application in deployment format
- create service, expose it
- enable HPA
- simulate load on deployment and verify HPA

-enable Metrics-Server

<https://github.com/kubernetes-sigs/metrics-server>

```
# kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

```
# kubectl edit deployments.apps --namespace kube-system metrics-server
```

```
- --kubelet-preferred-address-types=InternalIP
```

```
- --kubelet-insecure-tls
```

->modify the line, keep InternalIP

->append line

```
:wq!
```

```
# watch kubectl get nodes
```

```
# kubectl top nodes
```

-deploy an application in deployment format

```
# kubectl create deployment dpl1 --image nginx --replicas 2 -o yaml --dry-run=client >dpl1.yaml
```

```
# vim dpl1.yaml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  labels:
```

```
    app: dpl1
```

```
    name: dpl1
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      app: nginxdpl
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: nginxdpl
```

```
    spec:
```

```
      containers:
```

```
        - image: nginx
```

```
          name: nginxcnt
```

```
      resources:
```

```
        requests:
```

```
          cpu: "5m"
```

```
          memory: "8Mi"
```

```
        limits:
```

```
          cpu: "10m"
```

```
          memory: "12Mi"
```

```
:wq!
```

```
# kubectl create -f dpl1.yaml --dry-run=client
```

```
# kubectl create -f dpl1.yaml
```

```
# kubectl get deployments.apps
```

```
# kubectl get pods
```

```
# kubectl top pod -A
```

```
# kubectl top pod --namespace default
```

-create service, expose it

```
# kubectl expose deployment dpl1 --name dpl1ext --port 80 --protocol TCP --type NodePort
```

```
# kubectl get svc
```

```
dpl1ext NodePort 10.103.230.181 <none> 80:32254/TCP 11s
```

verify

open web browser

<http://192.168.29.104:32254/>

Welcome to nginx!

-enable HPA

```
# kubectl autoscale deployment dpl1 --max 6 --min 2 --cpu-percent 20 -o yaml --dry-run=client >dpl1hpa.yaml
# cat dpl1hpa.yaml
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  creationTimestamp: null
  name: dpl1
spec:
  maxReplicas: 6
  minReplicas: 2
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: dpl1
  targetCPUUtilizationPercentage: 20
status:
  currentReplicas: 0
  desiredReplicas: 0
:wq!
# kubectl autoscale deployment dpl1 --max 6 --min 2 --cpu-percent 20
# kubectl get hpa
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
dpl1 Deployment/dpl1 0%/20% 2 6 2 27s
# kubectl describe hpa
```

-simulate load on deployment and verify HPA

install **Siege** traffic emulator

<https://www.joedog.org/siege-home/>

<https://github.com/JoeDog/siege>

Siege is an open-source regression test and benchmark utility. It can stress test a single URL with a user defined number of simulated users, or it can read many URLs into memory and stress them simultaneously.

RHEL

-enable EPEL repository

<https://docs.fedoraproject.org/en-US/epel/>

-Install siege

yum install siege -y

<https://blog.eldernode.com/install-siege-on-centos/>

UBUNTU

apt-get install siege

siege --help

-Emulate traffic through siege for nginx deployment

duplicate MasterNode terminal:

terminal1

watch kubectl get pods

terminal2

watch kubectl top pods

terminal3, generate terrific

siege -q -c5 2m <http://192.168.29.104:32254/>

Now, check terminal1 and 2

number of pods will increase automatically

Now, stop siege with press Ctrl+c

and again, check terminal1

after around 3min number of pods will decrease to default.