

29Oct2022

Day17

Kubernetes Add, Remove, Rejoin Node to & from Cluster

Remove Node from k8S Cluster

kubectl get nodes

```
master1.example.com Ready control-plane 139d v1.24.1
node1.example.com Ready <none> 139d v1.24.1
node2.example.com Ready <none> 139d v1.24.1
node3.example.com Ready <none> 22m v1.24.1
node4.example.com Ready <none> 22m v1.24.1
node5.example.com Ready <none> 22m v1.24.1
```

kubectl create deployment dpl1 --image redis --replicas 7 -o yaml --dry-run=client >dpl1.yaml

ls

dpl1.yaml

kubectl create deployment dpl1 --image redis --replicas 7

kubectl get pods -o wide

Now, select node to remove. **node3.example.com**

before migration/cordon, create DaemonSet to check DaemonSet's behavior when a Node connect/disconnect to/from Cluster.

<https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

vim rsyslogdmnst.yaml

apiVersion: apps/v1

kind: DaemonSet

metadata:

name: rsyslogdmnst

namespace: kube-system

labels:

app: rsyslog

spec:

selector:

matchLabels:

name: rsyslogdmnst

template:

metadata:

labels:

name: rsyslogdmnst

spec:

containers:

- name: rsyslogcnt

image: rsyslog

:wq!

kubectl apply -f rsyslogdmnst.yaml --dry-run=client

kubectl apply -f rsyslogdmnst.yaml

kubectl get daemonsets.apps --namespace kube-system

| NAME | DESIRED | CURRENT | READY | UP-TO-DATE | AVAILABLE | NODE SELECTOR | AGE |
|--------------|---------|---------|-------|------------|-----------|---------------|-------|
| rsyslogdmnst | 5 | 5 | 5 | 5 | 5 | <none> | 5m52s |

Safely Drain a Node

<https://kubernetes.io/docs/tasks/administer-cluster/safely-drain-node/>

step1-disbale kube-scheduling on target node through 'cordon'

Cordon

kubectl get nodes

```
node3.example.com Ready <none> 29m v1.24.1
```

kubectl cordon node3.example.com

kubectl get nodes

```
node3.example.com Ready,SchedulingDisabled <none> 30m v1.24.1
```

uncordon

kubectl uncordon node3.example.com

kubectl get nodes

step2-Drain target node.

migrate workloads from target node to other nodes in cluster.

Drain

kubectl drain node3.example.com --force --ignore-daemonsets

node/node3.example.com already cordoned

WARNING: ignoring DaemonSet-managed Pods: kube-system/calico-node-n7jdc, kube-system/kube-proxy-k62mt, kube-system/rsyslogdmnst-tmg2r

evicting pod default/dpl1-58c79557cc-qd2wx

evicting pod default/dpl1-58c79557cc-n8279

pod/dpl1-58c79557cc-qd2wx evicted

pod/dpl1-58c79557cc-n8279 evicted

node/node3.example.com drained

kubectl get pods -o wide

step3-remove target node/nodes from Cluster

Delete

kubectl delete nodes node3.example.com

Add new Node to K8s Cluster

install new Linux Box.

OS: Ubuntu 18.04LTS

u: devops

p: ubuntu

After finish Linux Box installation, login through console and do basic configuration:

u: devops

p: ubuntu

\$ sudo -i

password: ubuntu

passwd

hostnamectl set-hostname node6.example.com

ip a s

route -n

cat /etc/rsyslog.conf

vim /etc/netplan/00-installer-config.yaml

This is the network config written by 'subiquity'

network:

version: 2

renderer: networkd

ethernets:

enp0s3:

dhcp4: false

addresses: [192.168.29.113/24]

gateway4: 192.168.29.1

nameservers:

addresses: [8.8.8.8,8.8.4.4,192.168.29.1]

:wq!

netplan apply

ping 8.8.8.8

Now, append node6 record in to '/etc/hosts' of all nodes in K8s Cluster

cat /etc/hosts

192.168.29.104 master1.example.com master1

192.168.29.110 master2.example.com master2

192.168.29.111 master3.example.com master3

192.168.29.105 node1.example.com node1

192.168.29.106 node2.example.com node2

192.168.29.107 node3.example.com node3

192.168.29.108 node4.example.com node4

192.168.29.109 node5.example.com node5

192.168.29.113 node6.example.com node6

192.168.29.112 haproxyb.example.com haproxyb

step1.prepare Node/Basic configuration

step2. install CRI, CNI, kubectl, kubeadm, kubelet

step3. check token on MasterNode

step4. join New Node to Cluster

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kube>

hostname

node6.example.com

sudo cat /sys/class/dmi/id/product_uuid

Container Runtimes

<https://kubernetes.io/docs/setup/production-environment/container-runtimes/>

Forwarding IPv4 and letting iptables see bridged traffic

cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf

overlay

br_netfilter

EOF

sudo modprobe overlay

sudo modprobe br_netfilter

#cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf

net.bridge.bridge-nf-call-iptables = 1

net.bridge.bridge-nf-call-ip6tables = 1

net.ipv4.ip_forward = 1

EOF

sudo sysctl --system

containerd

Installing containerd

Option 1: From the official binaries

<https://github.com/containerd/containerd/blob/main/docs/getting-started.md>

<https://github.com/containerd/containerd/releases>

wget <https://github.com/containerd/containerd/releases/download/v1.6.9/containerd-1.6.9-linux-amd64.tar.gz>

ls

containerd-1.6.9-linux-amd64.tar.gz

tar Cxvf /usr/local containerd-1.6.9-linux-amd64.tar.gz

systemd

```
# mkdir -p /usr/local/lib/systemd/system/
# vim /usr/local/lib/systemd/system/containerd.service
# Copyright The containerd Authors.
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.

[Unit]
Description=containerd container runtime
Documentation=https://containerd.io
After=network.target local-fs.target

[Service]
# Uncomment to enable the experimental sbservice (sandboxed) version of containerd/cni integration
#Environment="ENABLE_CRI_SANDBOXES=sandboxed"
ExecStartPre=-/sbin/modprobe overlay
ExecStart=/usr/local/bin/containerd

Type=notify
Delegate=yes
KillMode=process
Restart=always
RestartSec=5
# Having non-zero Limit*s causes performance problems due to accounting overhead
# in the kernel. We recommend using cgroups to do container-local accounting.
LimitNPROC=infinity
LimitCORE=infinity
LimitNOFILE=infinity
# Comment TasksMax if your systemd version does not supports it.
# Only systemd 226 and above support this version.
TasksMax=infinity
OOMScoreAdjust=-999

[Install]
WantedBy=multi-user.target
:wq!
# systemctl daemon-reload
# systemctl enable --now containerd
```

Installing runc

```
https://github.com/opencontainers/runc/releases
# wget https://github.com/opencontainers/runc/releases/download/v1.1.4/runc.amd64
# ls
containerd-1.6.9-linux-amd64.tar.gz runc.amd64
# install -m 755 runc.amd64 /usr/local/sbin/runc
```

Installing CNI plugins

```
https://github.com/containernetworking/plugins/releases
# wget https://github.com/containernetworking/plugins/releases/download/v1.1.1/cni-plugins-linux-amd64-v1.1.1.tgz
# ls
cni-plugins-linux-amd64-v1.1.1.tgz containerd-1.6.9-linux-amd64.tar.gz runc.amd64
# mkdir -p /opt/cni/bin
# tar Cxzf /opt/cni/bin cni-plugins-linux-amd64-v1.1.1.tgz
```

CRI-O

<https://github.com/cri-o/cri-o/blob/main/install.md#readme>

APT based operating systems

NOTE

set the environment variable \$OS as the appropriate OS as you installed

set the environment variable \$VERSION as the appropriate CRI-O version as you want installed

```
# echo "deb [signed-by=/usr/share/keyrings/libcontainers-archive-keyring.gpg]
https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/xUbuntu_18.04/ /" > /etc/apt/sources.list.d/devel:kubic:libcontainers:stable.list
```

```
# echo "deb [signed-by=/usr/share/keyrings/libcontainers-crio-archive-keyring.gpg]
https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable:/cri-o:/1.24/xUbuntu_18.04/ /" >
/etc/apt/sources.list.d/devel:kubic:libcontainers:stable:cri-o:1.24.list
```

```
# mkdir -p /usr/share/keyrings
```

```
# curl -L https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/xUbuntu_18.04/Release.key | gpg --dearmor -o
/usr/share/keyrings/libcontainers-archive-keyring.gpg
```

```
# curl -L https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable:/cri-o:/1.24/xUbuntu_18.04/Release.key | gpg --dearmor -o
/usr/share/keyrings/libcontainers-crio-archive-keyring.gpg
```

```
# apt-get update
```

```
# apt-get install cri-o cri-o-runc -y
```

```
# systemctl daemon-reload
```

```
# systemctl enable crio
```

```
# systemctl start crio
```

Docker Engine

<https://docs.docker.com/engine/install/ubuntu/>

```
# sudo apt-get update
```

```
# sudo apt-get install ca-certificates curl gnupg lsb-release
```

```
# sudo mkdir -p /etc/apt/keyrings
```

```
# curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg
```

```
# echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" |
sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

```
# sudo apt-get update
```

```
# sudo apt-get install docker-ce docker-ce-cli containerd.io docker-compose-plugin
```

cri-dockerd

<https://github.com/Mirantis/cri-dockerd>

<https://www.mirantis.com/blog/how-to-install-cri-dockerd-and-migrate-nodes-from-dockershim/>

```
# wget https://github.com/Mirantis/cri-dockerd/releases/download/v0.2.0/cri-dockerd-v0.2.0-linux-amd64.tar.gz
```

```
# tar xvf cri-dockerd-v0.2.0-linux-amd64.tar.gz
```

```
# sudo mv ./cri-dockerd /usr/local/bin/
```

```
# cri-dockerd --help
```

```
# wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-docker.service
```

```
# wget https://raw.githubusercontent.com/Mirantis/cri-dockerd/master/packaging/systemd/cri-docker.socket
```

```
# sudo mv cri-docker.socket cri-docker.service /etc/systemd/system/
```

```
# sudo sed -i -e 's,/usr/bin/cri-dockerd,/usr/local/bin/cri-dockerd,' /etc/systemd/system/cri-docker.service
```

```
# systemctl daemon-reload
```

```
# systemctl enable cri-docker.service
```

```
# systemctl enable --now cri-docker.socket
```

verify

```
# systemctl status cri-docker.socket
```

```
# systemctl status docker.service
```

```
# systemctl status containerd.service
```

```
# systemctl status crio.service
```

Installing kubeadm, kubelet and kubectl

```
# sudo apt-get update
```

```
# sudo apt-get install -y apt-transport-https ca-certificates curl
```

```
# sudo curl -fsSLo /usr/share/keyrings/kubernetes-archive-keyring.gpg https://packages.cloud.google.com/apt/doc/apt-key.gpg
```

```
# echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

```
# sudo apt-get update
```

```
# sudo apt-get install -y kubelet kubeadm kubectl
```

```
# sudo apt-mark hold kubelet kubeadm kubectl
```

```
# kubectl version --short
```

Join NEW node to existing K8s Cluster

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

Joining your nodes

to Join new node 'node6.example.com' new to write this command:

```
# kubeadm join --token <token> <control-plane-host>:<control-plane-port> --discovery-token-ca-cert-hash sha256:<hash>
```

Now, back to MasterNode and check/create token and ca.crt

```
# kubeadm token list
```

-to Create Token

```
# kubeadm token create
```

```
ludpbj.vqvvywe2q5qosl
```

-to create discovery-token-ca-cert-hash sha256

```
# openssl x509 -pubkey -in /etc/kubernetes/pki/ca.crt | openssl rsa -pubin -outform der 2>/dev/null | openssl dgst -sha256 -hex | sed 's/^.* //'
33db1c4ef050c21e798eddb51bf8a526b071f4f546be014faf0ef312a0546945
```

```
# hostname -i
```

```
192.168.29.104
```

Now, back to new node, 'node6.example.com'

```
# kubeadm join --token <token> <control-plane-host>:<control-plane-port> --discovery-token-ca-cert-hash sha256:<hash>
```

```
# swapoff -a
```

```
# free -m
```

```
# rm -rf /etc/containerd/config.toml
```

```
# systemctl restart containerd.service
```

```
# kubeadm join --token ludpbj.vqvvywe2q5qosl 192.168.29.104:6443 --discovery-token-ca-cert-hash
```

```
sha256:33db1c4ef050c21e798eddb51bf8a526b071f4f546be014faf0ef312a0546945 --cri-socket=/var/run/containerd/containerd.sock --v=5
```

NOTE

1-before join Debian-based OS Linux Box to existing/new K8s Cluster delete /etc/containerd/config.toml

```
root@node6:~# rm -rf /etc/containerd/config.toml
```

```
root@node6:~# systemctl restart containerd.service
```

```
root@node6:~# kubeadm join --token ludpbj.vqvvywe2q5qosl 192.168.29.104:6443 --discovery-token-ca-cert-hash
```

```
sha256:33db1c4ef050c21e798eddb51bf8a526b071f4f546be014faf0ef312a0546945 --cri-socket /var/run/containerd/containerd.sock --v=5
```

2-if in joining time any issue happens and procedure got failed, just delete Kubernetes directory from node and again try

```
root@node6:~# rm -rf /etc/kubernetes/
```

Rejoin Deleted Node from K8s Cluster again to K8s Cluster

after delete and maintenance, need rejoin **node3.example.com** to K8s Cluster.

on node3.example.com

```
# swapoff -a
```

```
# free -m
```

```
# rm -rf /etc/kubernetes/
```

```
# kubeadm join --token <token> <control-plane-host>:<control-plane-port> --discovery-token-ca-cert-hash sha256:<hash>
```

Now, back to MasterNode:

```
# kubectl get nodes
```

| NAME | STATUS | ROLES | AGE | VERSION |
|---------------------|--------|---------------|------|---------|
| master1.example.com | Ready | control-plane | 139d | v1.24.1 |
| node1.example.com | Ready | <none> | 139d | v1.24.1 |
| node2.example.com | Ready | <none> | 139d | v1.24.1 |
| node3.example.com | Ready | <none> | 2m1s | v1.24.1 |
| node4.example.com | Ready | <none> | 150m | v1.24.1 |
| node5.example.com | Ready | <none> | 150m | v1.24.1 |
| node6.example.com | Ready | <none> | 13m | v1.25.3 |