

31Oct2022

Day18

Kubernetes Cluster Upgrade

NetworkPolicy

Kubernetes Cluster Upgrade

<https://kubernetes.io/docs/tasks/administer-cluster/kubeadm/kubeadm-upgrade/>

The upgrade workflow at high level is the following:

- Upgrade a primary control plane node.
- Upgrade additional control plane nodes.
- Upgrade worker nodes.

```
# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master1.example.com	Ready	control-plane	141d	v1.24.1

```
# cat /etc/os-release
```

RHEL-Based

```
# cat /etc/redhat-release
```

Debian-Based

```
# lsb_release -a
```

Determine which version to upgrade to

```
# apt update
```

```
# apt-cache madison kubeadm
```

```
kubeadm | 1.25.3-00 | https://apt.kubernetes.io/kubernetes-xenial/main amd64 Packages ->target upgrade
```

Upgrading control plane nodes

The upgrade procedure on control plane nodes should be executed one node at a time.

```
# apt-mark unhold kubeadm && apt-get update && apt-get install -y kubeadm=1.25.3-00 && apt-mark hold kubeadm
```

```
# kubeadm version
```

```
# kubeadm upgrade plan
```

```
# sudo kubeadm upgrade apply v1.25.3
```

Drain the node

```
# kubectl drain <node-name> --ignore-daemonsets
```

```
# kubectl drain master1.example.com --ignore-daemonsets
```

Upgrade kubelet and kubectl

```
# apt-mark unhold kubelet kubectl && apt-get update && apt-get install -y kubelet=1.25.3-00 kubectl=1.25.3-00 && apt-mark hold kubelet kubectl
```

Restart the kubelet

```
# sudo systemctl daemon-reload
```

```
# sudo systemctl restart kubelet
```

Uncordon the node

```
# kubectl uncordon <node-name>
```

```
# kubectl uncordon master1.example.com
```

verify

```
# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master1.example.com	Ready	control-plane	141d	v1.25.3

Upgrade worker nodes

The upgrade procedure on worker nodes should be executed one node at a time or few nodes at a time, without compromising the minimum required capacity for running your workloads.

Upgrade kubeadm

```
root@master1:~# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master1.example.com	Ready	control-plane	141d	v1.25.3
node1.example.com	Ready	<none>	141d	v1.24.1

```
root@master1:~# ping node1.example.com
```

```
PING node1.example.com (192.168.29.105) 56(84) bytes of data.
```

```
64 bytes from node1.example.com (192.168.29.105): icmp_seq=1 ttl=64 time=0.324 ms
```

```
root@master1:~# ssh root@192.168.29.105
```

```
root@192.168.29.105's password: ubuntu
```

```
root@node1:~# apt-mark unhold kubeadm && apt-get update && apt-get install -y kubeadm=1.25.3-00 && apt-mark hold kubeadm
```

```
root@node1:~# sudo kubeadm upgrade node
```

```
root@node1:~# exit
```

Drain the node

```
root@master1:~# kubectl drain node1.example.com --ignore-daemonsets --force
```

```
root@master1:~# ssh root@192.168.29.105
```

```
root@192.168.29.105's password: ubuntu
```

Upgrade kubelet and kubectl

```
root@node1:~# apt-mark unhold kubelet kubectl && apt-get update && apt-get install -y kubelet=1.25.3-00 kubectl=1.25.3-00 && apt-mark hold kubelet kubectl
```

Restart the kubelet

```
root@node1:~# sudo systemctl daemon-reload
root@node1:~# sudo systemctl restart kubelet
root@node1:~# exit
```

Uncordon the node

```
root@master1:~# kubectl uncordon node1.example.com
```

verify

```
# kubectl get nodes
NAME                STATUS  ROLES    AGE  VERSION
master1.example.com Ready   control-plane  141d  v1.25.3
node1.example.com   Ready   <none>      141d  v1.25.3
```

NetworkPolicy

If you want to control traffic flow at the IP address or port level (OSI layer 3 or 4), then you might consider using Kubernetes **NetworkPolicies** for particular applications in your cluster.

The Two Sorts of Pod Isolation

There are two sorts of isolation for a pod:

isolation for egress	->outgoing traffic (from inside to outside)
isolation for ingress	->incoming Traffic (from outside to inside)

-By default, a pod is non-isolated for egress; all outbound connections are allowed

-By default, a pod is non-isolated for ingress; all inbound connections are allowed.

Implement NetworkPolicies

```
# kubectl get pods --namespace kube-system
# kubectl run pod1 --image nginx -o yaml --dry-run=client >pod1.yaml
# kubectl run pod2 --image centos -o yaml --dry-run=client >pod2.yaml
# kubectl run pod3 --image alpine -o yaml --dry-run=client >pod3.yaml
```

```
# vim pod1.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: pod1
  name: pod1
spec:
  containers:
  - image: nginx
    name: nginxcnt
:wq!
```

```
# vim pod2.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: pod2
  name: pod2
spec:
  containers:
  - image: centos
    name: centoscnt
    command: ['sleep','4800']
:wq!
```

```
# vim pod3.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: pod3
  name: pod3
spec:
  containers:
  - image: alpine
    name: alpinecnt
    command: ['sleep','4800']
:wq!
```

```
# kubectl apply -f pod1.yaml
# kubectl apply -f pod2.yaml
# kubectl apply -f pod3.yaml
# kubectl get pods
```

NetworkPolicies Scenarios

- 1-from same namespace(default)
- 2-from different namespaces

1-from same namespace(default)

config NetworkPolicies '**netpol1**' then **pod1** from **default** namespace will accept traffic from **pod3** in same namespace, not from other pods.

```
# kubectl config get-contexts
*      kubernetes-admin@kubernetes  kubernetes  kubernetes-admin

# kubectl get pods -o wide
pod1 1/1   Running  0      4m37s  172.16.11.67   node1.example.com      ->nginx
pod2 1/1   Running  0      4m35s  172.16.11.68   node1.example.com      ->centos
pod3 1/1   Running  0      4m32s  172.16.178.193 node3.example.com      ->alpine
verify connectivity
# kubectl exec -it pod2 -- ping 172.16.11.67
64 bytes from 172.16.11.67: icmp_seq=1 ttl=63 time=0.153 ms
# kubectl exec -it pod3 -- ping 172.16.11.67
64 bytes from 172.16.11.67: seq=0 ttl=62 time=0.418 ms
# kubectl exec -it pod2 -- curl 172.16.11.67
<title>Welcome to nginx!</title>

define NetworkPolicies
# kubectl api-resources | grep -i "network"
networkpolicies      netpol      networking.k8s.io/v1      true      NetworkPolicy
# vim netpol1.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: netpol1
  namespace: default
spec:
  podSelector:
    matchLabels:
      run: pod1                                ->TargetPod, who will receive requests from other SourcePods      Nginx(pod1)
  policyTypes:
    - Ingress
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              project: default
        - podSelector:
            matchLabels:
              run: pod3                                ->SourcePod/Pods, who will send requests to TargetPod      Alpine(pod3)

:wq!
# kubectl apply -f netpol1.yaml --dry-run=client
# kubectl apply -f netpol1.yaml
# kubectl get netpol
# kubectl describe netpol netpol1
verify
# kubectl exec -it pod2 -- ping 172.16.11.67
PING 172.16.11.67 (172.16.11.67) 56(84) bytes of data.
Ctrl+c
# kubectl exec -it pod2 -- curl 172.16.11.67
Ctrl+c
# kubectl exec -it pod3 -- ping 172.16.11.67
64 bytes from 172.16.11.67: seq=0 ttl=62 time=0.592 ms
```

2-from different namespaces

config NetworkPolicies '**netpol2**' then **pod1** from **default** namespace will accept traffic from **pod3** in **testsp** namespace, not from other pods from any namespaces.

```
# kubectl create namespace testsp
# kubectl get ns
# kubectl label namespaces testsp lbl=ns
# kubectl get namespaces testsp --show-labels
testsp  Active  4m29s  kubernetes.io/metadata.name=testsp,lbl=ns
# kubectl delete pod pod3 --force --grace-period=0
# kubectl get pods
pod1 1/1   Running  0      52m  172.16.11.67   node1.example.com
pod2 1/1   Running  0      52m  172.16.11.68   node1.example.com
# kubectl create -f pod3.yaml --namespace testsp
# kubectl get pods --namespace testsp
```

```
# vim netpol2.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: netpol2
  namespace: default          ->netpol2 namespace
spec:
  podSelector:
    matchLabels:
      run: pod1                ->TargetPod, that it will receive the traffic
  policyTypes:
    - Ingress
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            lbl: ns             ->SourcePod namespace label
      - podSelector:
          matchLabels:
            run: pod3           ->SourcePod name, that it will send traffic
:wq!
# kubectl create -f netpol2.yaml --dry-run=client
# kubectl create -f netpol2.yaml
kubectl get netpol
kubectl describe netpol netpol2
verify
# kubectl exec -it --namespace testsp pod3 -- ping 172.16.11.67
64 bytes from 172.16.11.67: seq=0 ttl=62 time=1.368 ms
# kubectl exec -it pod2 -- ping 172.16.11.67
PING 172.16.11.67 (172.16.11.67) 56(84) bytes of data.
Ctrl+c
```